

Scott Williams

COMPUTER SCIENCE PART II PROJECT DISSERTATION

**STEGANOGRAPHIC FILE SYSTEMS
WITHIN VIDEO FILES**

Christ's College
University of Cambridge

December 28, 2014

Performa

NAME:	Scott Williams
COLLEGE:	Christ's
PROJECT TITLE:	Steganographic file systems within video files
EXAMINATION:	Part II of the Computer Science Tripos
YEAR:	2015
WORD COUNT:	12,000
PROJECT ORIGINATOR:	Scott Williams
PROJECT SUPERVISOR:	Daniel Thomas

Original Aims of the Project

To investigate appropriate steganographic embedding methods for video and to develop a practical steganographic software package to enable the embedding of arbitrary data within video files via a file system interface. Raw AVI video files should be supported and a variety of steganographic embedding algorithms should be available. Basic file system commands should work within the presented logical volume and embedding should occur with no perceivable impact on video quality.

Work Completed

A complete software package has been developed enabling the embedding of arbitrary files within many video formats (including MP4 and AVI) via a file system interface. A total of 9 steganographic embedding algorithms are supported, along with encryption and plausible deniability functionality. Basic file system operations work as expected within the mounted volume and the embedding process can operate without any perceivable impact on video quality. Performance of the system is adequate for general use allowing high definition media content to be played directly out of the embedded volume.

Special Difficulties

None.

Declaration of Originality

I, Scott Williams of Christ's College, being a candidate for Part II of the Computer Science Tripos, hereby declare that this dissertation and the work described in it are my own work, unaided except as may be specified below, and that the dissertation does not contain material that has already been used to any substantial extent for a comparable purpose.

I give permission for my dissertation to be made available in the archive area of the Laboratory's website.

Signed:

Date:

Contents

1	Introduction	1
1.1	Motivation	1
2	Preparation	2
2.1	Background	2
2.1.1	Steganographic Concepts	2
2.1.2	Steganalysis	4
2.1.3	The AVI file format	5
2.1.4	JPEG compression	6
2.1.5	FFmpeg	7
2.1.6	Developing a file system	8
2.2	Existing tools	9
2.2.1	StegoStick	9
2.2.2	StegoMagic	10
2.2.3	TCSSteg	10
2.2.4	StegoVideo	10
2.2.5	OpenPuff	11
2.2.6	Steganosaurus	11
2.3	Requirements Analysis	11
2.3.1	Core Requirements	12
2.3.2	Possible Extensions	12
2.4	Choice of Languages and Tools	13
3	Implementation	14
3.1	Introduction	14
3.2	AVI Decoder	15
3.3	Steganographic Algorithms	18
3.3.1	LSB Sequential Embedding	19
3.3.2	Permuted LSB Embedding	23
3.3.3	Permuted and Encrypted LSB Embedding	25
3.4	File system	26
3.5	Extensions	26
3.5.1	Supporting multiple video formats	26
3.5.2	File system directory structures	26
3.5.3	Plausible deniability	26
3.5.4	Hiding data within audio streams	26
3.6	Testing	26
4	Evaluation	26
4.1	Satisfaction of Requirements	26

4.2	Correctness	26
4.3	Security	26
4.4	Performance	26
5	Conclusions	26
5.1	Future Project Directions	26

1 || Introduction

Steganography is the art of hiding information in apparently innocuous objects. Whereas cryptography seeks to protect only the content of information, steganography attempts to conceal the fact that the information even exists. This allows steganographic methods to be utilised in countries where encryption is illegal for example, or within the UK where keys for identified encrypted data can be forced to be handed over.

In this project I design and implement a practical steganographic software application - **Stegasis** - which enables users to embed arbitrary files within videos via a file system interface. **Stegasis** can operate with no perceivable impact on video quality and can achieve embedding capacities of upto 200% of the video size. A wide range of video formats are supported¹ along with several steganographic embedding algorithms. Standard encryption algorithms can be used to further protect embedded data and plausible deniability functionality protects users even when the presence of embedded data has been confirmed.

Steganographic methods operating on video have had comparatively little attention compared to images and audio. As such, there are few programs currently available which allow data to be steganographically hidden within video. **Stegasis** is the first application to enable the embedding of arbitrary files within videos via a file system interface.

1.1 Motivation

Digital media is ubiquitous on the Internet and high definition video content is now common place on video sharing and social networking websites. Video files of multiple gigabytes in size can reside on users devices without arousing suspicion, providing an ideal hiding place for large collections of sensitive files. Few programs are capitalising on this fact, and those that are, allow the user to embed only a single chosen file into a small range of video formats with very low embedding capacities. As with **TrueCrypt**², I believe that a practical system for protecting sensitive files should present the user with a mounted logical volume allowing the use of standard file system operations to create, access and organise embedded data. Furthermore, there exist many commonly used video formats along with many more currently in development. As such, a steganographic program operating on a small number of video formats not only greatly restricts usability, it will require constant development as new video formats inevitably become more popular. Instead, a generic solution applicable to a variety of video formats is preferred.

The many recent global surveillance disclosures show that using certain technologies related to privacy can get you “flagged” by authorities such as the NSA - it is no longer the case that simply encrypting data is enough to keep the owner safe.

¹Including many modern video formats such as MP4, MKV, FLV and AVI.

²A successful cryptographic program providing on-the-fly encryption and full disc encryption.

2 || Preparation

2.1 Background

In this section steganographic background material, definitions and concepts are introduced. A number of related technologies to the project are also discussed.

The most important property of any steganographic system is undetectability, that is, it should be impossible to differentiate between ordinary and steganographically modified objects. This requirement is famously formulated within Simmons' prisoners' problem.

Alice and Bob are imprisoned in separate cells and wish to formulate an escape plan. They are allowed to communicate, but all messages must pass through a warden Eve. If Eve suspects the prisoners of secretly discussing their escape plan, the communication channel will be severed and Alice and Bob thrown into solitary confinement. The prisoners attempt to utilise steganography to exchange details of their plan undetected. The steganographic system is considered broken if Eve is able to detect the presence of hidden messages within the prisoners exchanges. It is assumed that Eve has a complete knowledge of the steganographic algorithm being used, with the exception of the stego key, which Alice and Bob have agreed upon beforehand. This is in parallel with Kerckhoff's principle used within cryptography. The warden can be considered to be one of three categories: *passive*, *active* and *malicious*. A passive warden does not modify the exchanged messages in any way, whereas an active warden may modify the messages whilst maintaining their original meaning. For example an active warden may replace words with synonyms or reorder sentences. If images are being used as a transport medium then an active warden may recompress or crop the images. A malicious warden attempts to break the steganographic system and impersonate the prisoners in an attempt to obtain information.

This project is concerned with only the case of the *passive* warden. As such, any modification of the video files once `stegasis` has embedded data within them, will most likely render the embedded file system corrupt³.

2.1.1 Steganographic Concepts

A steganographic system consists of a number of individual components:

- A *Cover object* is the original object that the message will be embedded within.
- A *Message* is an arbitrary length sequence of symbols. For this project we consider a message $\mathcal{M} \in \{0, 1\}^n$ that is, a binary string.
- A *Stego key* is a secret key used within the embedding process.
- A *Stego object* is the result of embedding a message inside a cover object.

³This unfortunately means utilising video sharing websites such as YouTube and Facebook for distribution is not possible due to them performing compression upon video upload.

Definition 2.1. STEGANOGRAPHIC SYSTEM

Let \mathcal{C} be the set of all cover objects. For a given $\mathbf{c} \in \mathcal{C}$, let \mathcal{K}_c denote the set of all stego keys for \mathbf{c} , and the set \mathcal{M}_c denote all messages that can be communicated in \mathbf{c} . A steganographic system⁴, is then formally defined as a pair of embedding and extracting functions Emb and Ext ,

$$\begin{aligned} Emb : \mathcal{C} \times \mathcal{K} \times \mathcal{M} &\rightarrow \mathcal{C} \\ Ext : \mathcal{C} \times \mathcal{K} &\rightarrow \mathcal{M} \end{aligned}$$

satisfying,

$$\forall \mathbf{c}, \mathbf{k}, \mathbf{m}. \mathbf{c} \in \mathcal{C} \wedge \mathbf{k} \in \mathcal{K}_c \wedge \mathbf{m} \in \mathcal{M}_c \Rightarrow Ext(Emb(\mathbf{c}, \mathbf{k}, \mathbf{m}), \mathbf{k}) = \mathbf{m}$$

Definition 2.2. EMBEDDING CAPACITY

The Embedding Capacity (payload) \mathcal{P}_c for a given cover object $\mathbf{c} \in \mathcal{C}$ is defined in bits as,

$$\mathcal{P}_c = \log_2 |\mathcal{M}(\mathbf{c})|$$

The relative embedding capacity \mathcal{R}_c for a given cover object $\mathbf{c} \in \mathcal{C}$ is defined as,

$$\mathcal{R}_c = \frac{\log_2 |\mathcal{M}(\mathbf{c})|}{n}$$

where n is the number of elements in \mathbf{c} .

For example, consider \mathcal{C} to be the set of all 512×512 greyscale images, embedding one bit per pixel gives $\mathcal{M} = \{0, 1\}^{512 \times 512}$ and $\forall \mathbf{c} \in \mathcal{C}. |\mathcal{M}(\mathbf{c})| = 2^{512 \times 512}$. The embedding capacity $\forall \mathbf{c} \in \mathcal{C}$ is then $512 \times 512 \approx 33\text{kB}$ as expected. In this case, n is equal to the number of pixels in \mathbf{c} and therefore the relative embedding capacity is equal to 1 bpp (bits per pixel), again as expected.

Using the definitions above, we can define a simple expression for the embedding capacity of a video file.

Definition 2.3. EMBEDDING CAPACITY FOR VIDEO

With \mathcal{C} as the set of all video files, the embedding capacity \mathcal{V}_c for a given video $\mathbf{c} \in \mathcal{C}$ can be expressed as,

$$\mathcal{V}_c = \sum_{f \in frames(\mathbf{c})} \mathcal{P}_f$$

Note that for certain embedding algorithms, the embedding capacity can depend on both the input data and the cover object⁵. However, in some cases the following expression is also valid,

$$\mathcal{V}_c = |frames(\mathbf{c})| \cdot \mathcal{P}_{f_0}$$

⁴This is specifically steganography by cover modification.

⁵Many algorithms operating on JPEG images for example will not embed within zero valued DCT coefficients.

Definition 2.4. STEGANOGRAPHIC CAPACITY

The concept of Steganographic Capacity is loosely defined as the maximum number of bits that can be embedded within a given cover object without introducing statistically detectable artifacts.

For completeness, the least significant bit (LSB) of a given number is defined as follows,

$$\text{LSB}(x) = x \bmod 2$$

It will be useful to visually inspect the effect of steganographic embedding algorithms operating on the LSBs of pixels. The *LSB Plane* of an image is therefore defined.

Definition 2.5. LSB PLANE

The Least Significant Bit Plane of a given image \mathbf{c} and a specified colour channel q is defined as the 1 bit image $\text{LSBP}(\mathbf{c}, q)$ which has resolution equal to that of image \mathbf{c} and with pixel values $\text{LSBP}(\mathbf{c}, q)(x, y)$ given by,

$$\text{LSBP}(\mathbf{c}, q)(x, y) = \text{LSB}(\mathbf{c}(x, y))$$

2.1.2 Steganalysis

Steganalysis is the study of detecting messages embedded using steganographic techniques; this is analogous to cryptanalysis applied to cryptography. A steganalysis attack is considered successful (that is, the steganography has been broken) if it is possible to correctly distinguish between cover and stego objects with probability better than random guessing. Note that it is not necessary to be able to read the contents of the secret message to break a steganographic system.

A trivial example of steganalysis arises when the steganalyst has access to the original cover object used within the embedding procedure. By computing the difference between the stego and cover objects, the steganalyst can immediately detect the presence of a hidden message. This attack identifies a number of important points to consider when developing a practical steganographic system. Firstly, embedding within popular media content should be discouraged, as the cover object will be likely widely available. Secondly, if a user is embedding within original content, for example a video recorded by them, any copies of the original file should be securely erased after embedding.

Steganalysis methods can be split into two main categories, *Targeted Steganalysis* and *Blind Steganalysis*. Targeted Steganalysis occurs when the steganalyst has access to the details of the steganographic algorithm used for embedding. The steganalyst can accordingly target their activity to the specific stegosystem. On the other hand, if the steganalyst has no knowledge of the utilised steganographic algorithm, Blind Steganalysis techniques must be applied. In this project, Targeted Steganalysis attacks are developed for several of the proposed embedding algorithms.

2.1.3 The AVI file format

As specified within the project proposal, this project only (initially) looks at raw uncompressed AVI files. Furthermore, only AVI version 1.0⁶ files are investigated and therefore supported natively⁷ by **Stegasis**. Unfortunately, uncompressed AVI is today, a very uncommon video format. This is likely due to its relatively huge file sizes when compared to a modern compressed format such as MP4 H.264. For example, one minute of 720p HD footage encoded as uncompressed AVI is roughly 4.2 GB.

The AVI file format is a Resource Interchange File Format (RIFF) file specification developed by Microsoft and originally introduced in November 1992. The data within RIFF files is divided into chunks and lists, each of which is identified by a FourCC tag. An AVI file takes the form of a single chunk in a RIFF formatted file, which is then subdivided into two mandatory lists and one optional chunk. The first sub-list is the file header containing metadata about the video (for example framerate, width and height). The second sub-list contains the actual audio/video data and the optional chunk indexes the offsets of the data chunks within the file.

We therefore have an AVI file laid out as follows, see the appendix for a more detailed expanded form.

```
RIFF ( 'AVI' 
    LIST ( 'hdrl' ... )
    LIST ( 'movi' ... )
    [ 'idx1' (<AVI Index>)]
)
```

Listing 2.1: AVI RIFF form

With a RIFF chunk being defined as follows:

```
struct CHUNK {
    char fourCC[4],
    int ckSize,
    char ckData[ckSize] // contains headers or video/audio data
};
```

Listing 2.2: RIFF chunk

And a RIFF list defined as:

```
struct LIST {
    char listCC[4], // Will always be the literal 'LIST'
    int listSize,
    char listType[4],
    char listData[listSize]
};
```

Listing 2.3: RIFF list

⁶Not including the Open-DML extension (version 1.02).

⁷All other video formats (including compressed AVI) are supported via the use of FFmpeg, as described in section 3.5.1.

An AVI file consists of a number of data streams (usually 2, one for audio and one for video) interleaved within the movi list. Each stream is identified by a FourCC tag consisting of a two-digit stream number followed by a two-character code listed in table 2.1.

Two-character code	Description
db	Uncompressed video frame
dc	Compressed video frame
pc	Palette change
wb	Audio data

Table 2.1: AVI stream types

Each stream has a corresponding AVI stream header and format chunk within the above mentioned hdrl list. These data structures contain information about the stream including the codec and compression used (if any). Specifically, the `fccHandler` field contains a FourCC tag that identifies a specific data handler. For raw uncompressed video this will equal ‘DIB’ (Device Independent Bitmap). Any user provided AVI files with a `fccHandler` not equal to ‘DIB’, that is, the AVI contains compressed video, will at this point be rejected and an error message presented to the user.

The movi list contains the raw video and audio data within sequential RIFF chunks. Each chunk for the DIB video stream contains one frames worth of pixel data, with each pixel represented by a 3 byte BGR (Blue Green Red) triple - a total of 24 bits per pixel. The first 3 byte triple corresponds to the lower left pixel of the final image⁸.

If we use an embedding algorithm which embeds 3 bits per pixel (that is, 1 bit per colour channel per pixel) we can derive an expression for the embedding capacity of a video \mathbf{c} in terms of the height h and width w in pixels, the total number of frames t and the frame rate f in frames per second:

$$\mathcal{V}_\mathbf{c} = \frac{3 \cdot w \cdot h \cdot t}{f}$$

These values are all available within the `AVIMAINHEADER` structure allowing the user to be informed of the video’s embedding capacity upon formatting.

2.1.4 JPEG compression

The JPEG file format will prove useful when developing a universal steganographic technique operating across many video formats, see section ???. Steganography within JPEGs has had a comparatively large amount of attention from the research community, most likely due to their popularity and the fact that virtually every camera will produce images in the JPEG format. As such, there exists a fair number of well documented steganographic embedding algorithms for JPEG.

The JPEG compression process consists of 5 main procedures:

⁸This can be inverted via the use of an option within the `BITMAPINFOHEADER`.

1. Transform the image into an optimal color space.
2. Downsample chrominance components by averaging groups of pixels together.
3. Apply a Discrete Cosine Transform (DCT) to blocks of pixels.
4. Quantise each block of DCT coefficients using a quantisation table.
5. Encode the resulting coefficients using a Huffman variable word-length algorithm.

Note that step 4 is an example of lossy compression, whereas step 5 is lossless. Therefore most steganographic algorithms will operate on the quantised DCT coefficients (between steps 4 and 5) to avoid embedded data being lost.

Conveniently, the Independent JPEG Group provide the `libjpeg` C library which which will abstract the complexities of the JPEG format and allow direct access to the quantised DCT coefficients prior to step 5 being executed.

It is worth noting that the JPEG decompression and compression processes are computationally expensive. This is especially important when dealing with video since the average 3 minute music video, for example, consists of 4,500 frames (which can be considered as individual JPEGs). Since performance of the virtual file system is important, design decisions will need to be made to accommodate this. Also worth noting is that although JPEG files are small on disk, they're not once decompressed into RAM. It will not be possible to hold all 4,500 decompressed JPEG frames of the average music video in RAM, which is unfortunate again for performance reasons.

2.1.5 FFmpeg

`FFmpeg` is an open source, multimedia framework. It is a “complete, cross-platform solution to record, convert and stream audio and video”. In particular, it contains codecs for nearly every video format available today.

The pitfalls of the uncompressed AVI video format, as discussed in section 2.1.3, show that `Stegasis` would greatly benefit from operating on multiple video formats other than uncompressed AVI. I could continue to investigate more video formats and develop codecs for these as part of the project. However, this will become a very time consuming endeavor most likely resulting in very brittle, untested parsers. Instead, it would be wise to leverage the `FFmpeg` framework for this functionality.

One trivial solution to allow `Stegasis` to operate on multiple video formats would be to convert all user provided video files to uncompressed AVI, using `FFmpeg`, prior to the embedding process. However, this doesn't solve the problems of the huge file sizes and uncommonality of the format. Note that it is not possible to convert to uncompressed AVI, perform the embedding and then convert back to the original provided format since the conversion process will be lossy, damaging the embedded data.

A novel solution to this problem is posed in section ?? and makes use of `FFmpeg` for the video conversion.

2.1.6 Developing a file system

A file system can either operate within *kernel* or *user space*. It was decided at the project proposal stage to develop the file system component for **Stegasis** in user space using the FUSE (Filesystem in Userspace) library for a number of reasons. Firstly, developing a kernel module is complex and hard to test - a **Segmentation fault** occurring within kernel space code will bring down the entire machine. A kernel module also requires a large amount of boiler plate code and I would prefer to spend time on the steganographic portion of this project rather than getting bogged down with the complexities of a kernel file system implementation. In contrast, FUSE ships with an example “hello world” file system which is less than 100 lines of C code. Secondly, developing the file system in user space will cause the final application to be a lot more portable and easier for users to install - a kernel space file system would require super user permission to load the related kernel module.

There are however disadvantages to using a file system in user space, performance being one of them. This is due to the FUSE kernel module having to act as a proxy between the system call and the user space code. This is explained below.

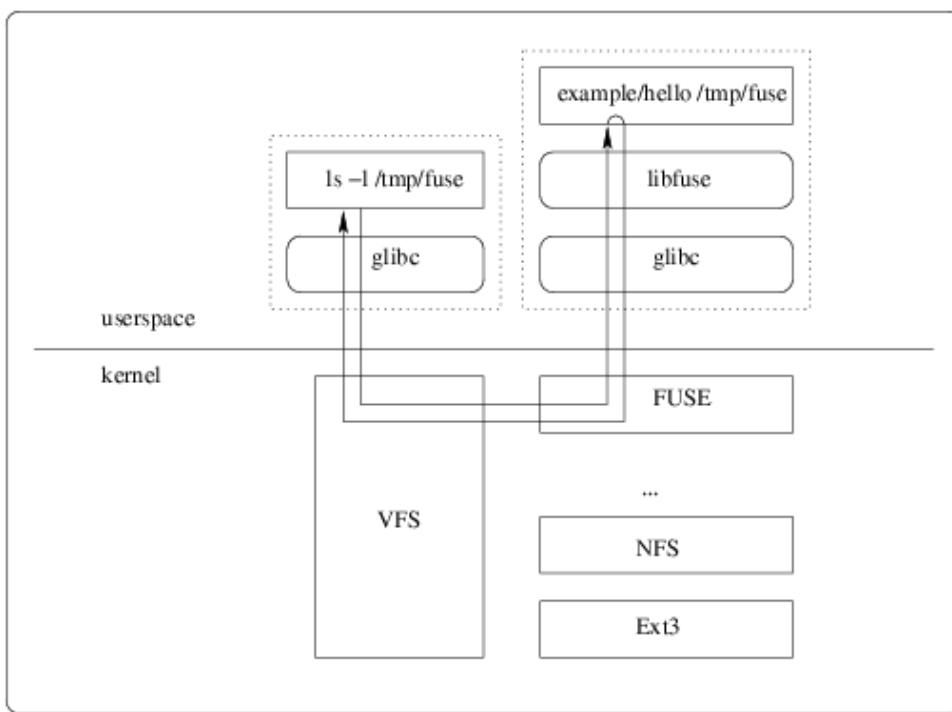


Figure 2.1: Path of a file system call in the hello world example.

Figure 2.1 shows the path of a file system call in the provided hello world example file system. We can see the FUSE kernel module acting as a proxy between the VFS system call and the `example/hello` user space code. A kernel space file system would not need

to re-enter user space to complete the system call, hence giving better performance.

The FUSE library provides a number of function definitions which the user space code implements. These functions are then called when the corresponding file system operation occurs. Some of the important operations are discussed in detail below.

```
int read(const char *path, char *buf, size_t size, off_t offset,  
        struct fuse_file_info *fi);
```

Listing 2.4: FUSE read operation.

The `read` function is called when a file system read occurs. It requests that `size` bytes of the file `path` starting at offset `offset` should be written to the buffer `buf`.

The `write` function is similar:

```
int write(const char *path, const char *buf, size_t size, off_t  
         offset, struct fuse_file_info *fi);
```

Listing 2.5: FUSE write operation.

It requests that `size` bytes from the buffer `buf` should be written to the file `path` starting at offset `offset`.

Similar functions exist for all of the standard file system operations, see appendix section ??? for a full list.

2.2 Existing tools

The relatively little work on steganography within video is reflected in my search for steganographic programs operating on video files. This section contains an exhaustive list of all the video steganography tools I could find freely⁹ available on the Internet. A total of 6 tools claimed to provide steganographic embedding functionality within video files. Of these 6, only 3 actually attempt to embed within the video data itself. None of the identified programs allow the user to embed more than one file¹⁰ and none of them provide any sort of file system interface.

2.2.1 StegoStick

StegoStick claims to allow users to “hide any file into any file”. This statement suggests that the program is simply appending the requested file to the end of the cover object. This suspicion is partly true; based on the file extension, **StegoStick** splits cover objects into 3 categories: images, media and other. The other category does indeed just append the file to the cover object, whereas the image and media category do attempt to employ steganographic embedding methods. The images category applies to files with extensions

⁹A further 2 programs exist claiming to embed within video, however these are closed source and not freely available to download. Therefore they have been excluded from this list. (Info Stego, Hiderman)

¹⁰Admittedly you could embed a compressed archive using these tools to effectively allow a directory structure to be embedded.

JPG, GIF and BMP and uses LSB embedding within BMP files (other image formats are converted to BMP prior to embedding). The media category applies to WAV, AVI and MPG files and assumes each format has a “header” of 44+55 bytes¹¹. Although this seems to be true for the WAV format, this is not the case for AVI nor MPG files. StegoStick will then use blind LSB embedding within the remaining data. As such, my attempts to use StegoStick to embed within AVI files rendered the resulting video unplayable.

2.2.2 StegoMagic

StegoMagic claims to “work on all types of files and all size of data” which again sounds as though it’s appending the file to the end of the cover object. This is indeed the case, embedding an image within a video and inspecting the modified file shows that data has just been appended to the end of the video, albeit encrypted. StegoMagic does not specify the encryption algorithm used and the source code is not available to view. Furthermore, the user cannot specify an encryption key to use. Instead, StegoMagic generates a 5 digit number during the embedding process and echos this to the user.

2.2.3 TCSteg

TCSteg is a Python script accompanying a blog post written by Martin Fiedler discussing hiding TrueCrypt volumes within MP4 files. The method described embeds the TrueCrypt volume within the MP4 atom `mdat` and modifies the chunk offset table within the `moov` atom so that any application playing the video will ignore the embedded data. A nice property of TCSteg is that the resulting video file can be directly mounted by TrueCrypt since it ignores the MP4 header data prior to the embedded volume.

The above programs all resort to embedding within video files by either appending the embedded data to the end of the video, or inserting the embedded data at some point within the video file. I do not consider this approach to embedding data secure, and it should be a trivial task for any steganalyst to detect the presence of embedded data within the stego objects using a simple hex editor. Therefore, the above stegosystems should be considered broken and definitely not used for the hiding of sensitive data.

2.2.4 StegoVideo

StegoVideo is a Virtual Dub filter¹² which allows users to embedded a file within AVI files (supporting multiple compression codecs). I am unsure of the exact steganographic embedding algorithm used since the program is closed source, but the website does mention that StegoVideo makes use of error correction codes to allow embedded data to be recovered even after the resulting video has been compressed - although this is understandably dependant on the compression amount. StegoVideo claims to protect the embedded data via the use of a passkey (a 5 digit number), although as with StegoMagic,

¹¹Listed in the source as “44 byte header + 54 bytes of extension space”.

¹²Which is also available in a stand alone executable form.

this is not provided by the user and is instead generated and presented to the user to make a note of.

2.2.5 OpenPuff

`OpenPuff` is a steganographic tool supporting a wide range of formats, including 3GP, MP4, MPG and VOB. It allows users to embed a file within a collection of carrier objects and uses 3 user provided passwords to encrypt, scramble and whiten (mixing with a high amount of noise) the provided file. Plausible deniability is also provided via the option to add decoy content. `OpenPuff` successfully embedded and retrieved a text file within a sample MP4 video and I could notice no perceivable impact on video quality. Performance was also good due to multithreading support. However, the embedding capacity is very limited. A hard limit of 256 MB is imposed regardless of the number and size of the carrier objects and I was only able to achieve embedding capacities of around 0.0043%¹³ even at the maximum capacity setting. This makes `OpenPuff` impractical for hiding large files - for example, you would need around 770 60 MB MP4 carrier files to embed a standard 2 MB JPEG image.

2.2.6 Steganosaurus

`Steganosaurus` is a cross platform steganographic program developed by James Ridgway. It allows users to embed a file within MP4 (H264) files via the modification of motion vectors. The input file is encrypted using AES with a user provided pass phrase. Need to test this in Linux.

The above 3 programs are much more promising from a steganographic security point of view and some of them also support multiple video formats. However, all feature the same limitation of only allowing the user to embed one chosen file and the offered embedding capacities are far from practical for use with large files.

This project aims to remedy these issues by providing the user the opportunity to embed an arbitrary number of files within a video via a file system interface and providing high capacity steganographic embedding algorithms offering capacities in excess of 100%¹⁴ of the cover object's size¹⁵.

2.3 Requirements Analysis

After reviewing the necessary background material and investigating current available solutions to the problem of steganography within video, the following collection of requirements were produced. For the project to be considered a success, at least all of the core requirements should be fulfilled.

¹³2,600 bytes within a 60 MB video.

¹⁴This specifically means a file of size n bytes can have more than n hidden bytes residing inside it.

¹⁵This is very much a trade off - capacities larger than the file size will come at the sacrifice of steganographic security. However, this decision is presented to the user rather than decided by the program itself.

2.3.1 Core Requirements

Stegasis should:

- Allow users to embed data within video files:
 - Several steganographic embedding algorithms should be available.
 - Each embedding algorithm, \mathcal{A} , should satisfy correctness. That is,
$$\forall \mathbf{c}, \mathbf{k}, \mathbf{m}. \text{Ext}_{\mathcal{A}}(\text{Emb}_{\mathcal{A}}(\mathbf{c}, \mathbf{k}, \mathbf{m}), \mathbf{k}) = \mathbf{m}.$$
 - Steganalysis tools should be developed to test the security of the proposed embedding algorithms.
 - An optional user provided password should encrypt data prior to embedding.
 - A capacity option should allow users to specify the percentage of each video frame to embed within.
- Provide a file system interface:
 - The presented logical volume should reside at a user provided mount point.
 - Data written to the file system should be embedded on the fly within the chosen video file.
 - Data accessed from the file system should be retrieved on the fly from within the video.
 - Standard file system operations such as creating, deleting and moving files should work as expected, and standard Unix tools such as `cp`, `mv` and `rm` should also work as expected.
- Support raw uncompressed AVI video:
 - Uncompressed AVIs should be natively parsed allowing access to individual pixel data.
- Provide performance adequate for normal use:
 - Full HD video content should be playable directly from within the presented file system.

2.3.2 Possible Extensions

If time constraints allow, the following extension tasks shall also be completed.

Stegasis should:

- Support a wide range of video formats:
 - Specifically including the popular video format MP4.
- Allow directory operations within the file system:

- Creating directories using the `mkdir` command should work as expected, as should using the `mv` and `rm` commands.
 - Organising files within directories should also work as expected.
- Embed also within audio data:
 - Data should also be embedded within the (possible) audio stream of the video, therefore increasing the embedding capacity.
- Provide plausible deniability:
 - A second file system should be (optionally) embedded within the video, mountable with a second passphrase.
 - The presence of the second, hidden file system should not be detectable.
- Be evaluated for perceivable video impact using a web application:
 - The web application should evaluate the claim “Embedding has no perceivable impact on video quality.” by obtaining data from multiple users.

2.4 Choice of Languages and Tools

With the above requirements for the final product defined, an appropriate set of languages and tools can be identified.

It is first noted that **Stegasis** (as developed for this project) will only function on the **Linux** operating system. That is, there is no requirement for **Stegasis** to be cross platform.

The file system is an important aspect of **Stegasis** and so it is initially decided which approach to take in developing it as this will influence the later choice of an appropriate programming language. As described within section 2.1.6, it was decided to use the **FUSE** library to develop the file system component in user space.

We now address the choice of primary programming language for the development of **Stegasis**.

Several of the core (and extension) requirements strongly suggest a lower level language such as C or C++ rather than a higher level sandboxed language such as Java. For example, the parsing and modification of AVI files lends itself to a language like C since it will involve large amounts of byte level manipulation. Furthermore, the Microsoft file format reference defines the different data structures used within AVIs as C **structs**. The identified library for implementing the file system aspect - **FUSE** - is natively a C library (as is the **libjpeg** library, and libraries provided by **FFmpeg**). Although wrappers for other languages (including Java) do exist, they seem to be lacking documentation and few are being actively maintained. The requirement that **Stegasis** should support several steganographic embedding algorithms implores the use of object oriented techniques; defining a **SteganographicAlgorithm** interface of which each embedding algorithm

implements. This suggests C++ over C. The final core requirement, performance, also favours C/C++ over Java¹⁶ due to the JVM overheads.

The reasons above led to the conclusion that C++ should be the primary language used to develop **Stegasis**.

As discussed in section 2.1.5, FFmpeg will be used for the extension task “**Stegasis** should support a wide range of video formats”, to allow the decoding and conversion of the many video formats available today, together with library libjpeg discussed in section 2.1.4 for the manipulation of JPEG images.

During the implementation of **Stegasis**, a number of small steganalysis programs will be developed. These will likely be written in a scripting language such as Python or Matlab since both have extensive library support for mathematical operations.

The extension task “**Stegasis** should be evaluated for perceivable video impact using a web application” will require a website to be developed and hosted for easy access to participants and a database to store the collected user data. Node.js together with the web application framework Express and the database MongoDB was chosen as the development stack for the site. This decision was mainly due to the speed at which you can develop *CRUD* (create, read, update and delete) web applications - essentially what this evaluation site is - and my previous experience with the technologies.

3 || Implementation

3.1 Introduction

The development of **Stegasis** consisted of the 5 main stages sectioned within this chapter. Firstly, a parser for the AVI file format as discussed in section 2.1.3 was developed allowing direct access to video pixel data. Next, steganographic embedding algorithms were implemented along with corresponding steganalysis tools to test the security of the proposed techniques. The file system was then developed utilising the AVI decoder and steganographic algorithms to embed and extract data directly into and out of video files. Finally, the extension tasks were individually addressed providing support for multiple video formats, directory structures and plausible deniability¹⁷. The testing section provides an overview of the testing processes applied throughout development.

The actual software development process taken differs from that laid out below; each section was not wholly completed before moving onto the next. Instead, an iterative process was taken across all sections, embracing the modern “Launch early, iterate often” methodology. For example, as specified in the project proposal timetable, a simplified version of **Stegasis** was initially produced only offering one simple embedding algorithm and basic file system functionality. This allowed integration issues to be identified early on, when the code was still very malleable. Once this basic version was working, an

¹⁶There have been numerous studies showing that C/C++ code performs better than equivalent Java code.

¹⁷The evaluation site extension task is discussed within the evaluation chapter.

iterative approach was then taken to add more functionality and features. For the sake of readability, I have structured the sections below to group together implementation details for each separate concept.

3.2 AVI Decoder

The concept of an **AVI** decoder is first abstracted to that of a generic **Video Decoder** interface¹⁸. The core requirements state that the **AVI** decoder should allow access to individual pixel data. The pixel data within an **AVI** file is grouped into chunks, one per video frame. It was therefore decided to define the **Video Decoder** to allow access to the video pixel data at a granularity of a single video frame. It would also be useful for the **Video Decoder** interface to expose metadata about the video, for example, the total number of video frames in the video, the height and width of the video frames and the total size (in bytes) of each video frame.

This gives the following definition for the **Video Decoder** interface (`NextFrameOffset` will be discussed in section 3.4):

```
class VideoDecoder {
public:
    virtual Chunk *getFrame( int frame ) = 0;
    virtual int getFileSize() = 0;
    virtual int getNumberOfFrames() = 0;
    virtual int getFrameSize() = 0;
    virtual int getFrameHeight() = 0;
    virtual int getFrameWidth() = 0;

    virtual void getNextFrameOffset( int *frame , int *offset ) = 0;
    virtual void setNextFrameOffset( int frame , int offset ) = 0;

    virtual void setCapacity( char capacity ) = 0;
    virtual void writeBack() = 0;
    virtual ~VideoDecoder() {};
};
```

Listing 3.1: Video Decoder interface (`video/video_decoder.h:15`)

Note that `getFrame` returns a `Chunk` wrapper object, rather than a raw `char` pointer to the frame pixel data, adhering to the *Dependency Inversion* principle. This will be useful when dealing with different video formats that don't necessarily group all of a frames video data to be accessible by a single `char` pointer.

A `Chunk` abstracts the concept of a single frames video data. In the case of uncompressed **AVI**, this can be thought of as a `char` pointer to the GBR pixel data, along with an associated frame size in bytes. A boolean value is also associated with each chunk, signifying if the chunk data has been modified, that is, it is dirty.

The `Chunk` interface is therefore defined as follows:

¹⁸The term “interface” is used as shorthand for an abstract C++ base class. That is, a class with pure virtual member functions and no function implementations.

```

class Chunk {
    protected:
        long chunkSize;
    public:
        virtual long getChunkSize() = 0;
        virtual char *getFrameData(int n=0, int c=0) = 0;

        virtual bool isDirty() = 0;
        virtual void setDirty() = 0;
};

```

Listing 3.2: Chunk interface (`video/video_decoder.h:4`)

Note that the parameters for `getFrameData` are optional. For the AVI decoder, these will not be used.

The AVI parssing process can be thought of consisting of two main parts; parsing the video headers and parsing the video chunk data. The following pseudocode illustrates this with the headers being parsed lines 1-12 and the chunks being parsed lines 15-22:

Algorithm 3.1 AVI parsing process

```

1: f ← open(file_path)
2: riff_header ← readRiffHeader(f)
3: if riff_header.fourCC != RIFF then
4:     print “File is not an AVI file”
5:     Exit
6: avi_header ← readAviHeader(f)
7: bitmap_info_header ← readBitmapInfoHeader(f)
8: if bitmap_info_header.compression != 0 then
9:     print “Stegasis does not natively support compressed AVI files”
10:    print “Rerun using the -f flag”
11:    Exit
12: audio_info_header ← readAudioInfoHeader(f)
13: frame_chunks ← []
14: i ← 0      ▷ File pointer is now positioned at the start of the audio video chunks
15: while i < avi_header.total_frames do
16:     chunk ← readChunk(f)
17:     if chunk.fourCC == 00db then
18:         frame_chunks[i].chunkSize = chunk.chunkSize
19:         frame_chunks[i].frameData = readChunkData(f)
20:         i ++
21:     else
22:         Advance f chunk.chunkSize bytes           ▷ Chunk was not a video chunk

```

See the appendix section B for some longer code samples. The actual implementation is slightly more complex than presented above due to the existence of JUNK chunks. The AVI file format specifies that any number of chunks with a FourCC code of JUNK and

of arbitrary length can be inserted between any AVI list structures. The parser must therefore be able to cope with this.

The `WriteBack` function of the AVI decoder will write back any modified Chunk data into the original AVI file. This operation is described in algorithm 3.2 below.

Algorithm 3.2 AVI write back process

```

1: Seek f to the chunks offset
2: i ← 0
3: while i < avi_header.total_frames do
4:   chunk ← readChunk(f)
5:   if chunk.fourCC == 00db then
6:     if frame_chunks[i].isDirty then
7:       Write frame_chunks[i].frameData to f
8:       frame_chunks[i].dirty = false           ▷ This chunk is no longer dirty
9:     else
10:      Advance f chunk.chunkSize bytes    ▷ Chunk did not need to be written
11:      i ++
12:    else
13:      Advance f chunk.chunkSize bytes    ▷ Chunk was not a video chunk

```

The remaining functionality of the parser is mostly trivially returning data from the `aviHeader` structure (for example `aviHeader.height`, `aviHeader.totalFrames` etc.) The only other function of interest is `getFrameSize`, which returns the number of bytes within each frame that can be embedded within. This (incorrectly) assumes that each frame can always have the same number of bytes embedded within it. This assumption is however true for the steganographic algorithms presented in the next section¹⁹, and is implemented as in listing 3.4.

```

virtual int frameSize() {
    return (int)floor(this->aviHeader.width * this->aviHeader.height
                     * 3 * (capacity / 100.0));
}

```

Listing 3.3: AVI decoder `frameSize` function (`video/avi_decoder.cc:298`)

This expression arises from the fact that uncompressed AVI uses 24 bits (3 bytes) per pixel value. So since there are $height \cdot width$ pixels within a single frame, we simply multiply this by 3 to get the total number of bytes. `Capacity` is a user provided percentage ranging in value from 1 - 100. It specifies the percentage of the frame to embed within. `frameSize` must therefore reduce the returned frame size value by “capacity percent”.

The effect of the capacity parameter is illustrated within figure 3.1. The top left image is the original video frame and the top right image is the LSB plane (red channel) of the frame with no data embedded. The bottom two images have data sequentially embedded within them using capacity settings of 50% and 15% respectively.

¹⁹It is also true for the simple JPEG embedding algorithms used within the extension tasks.

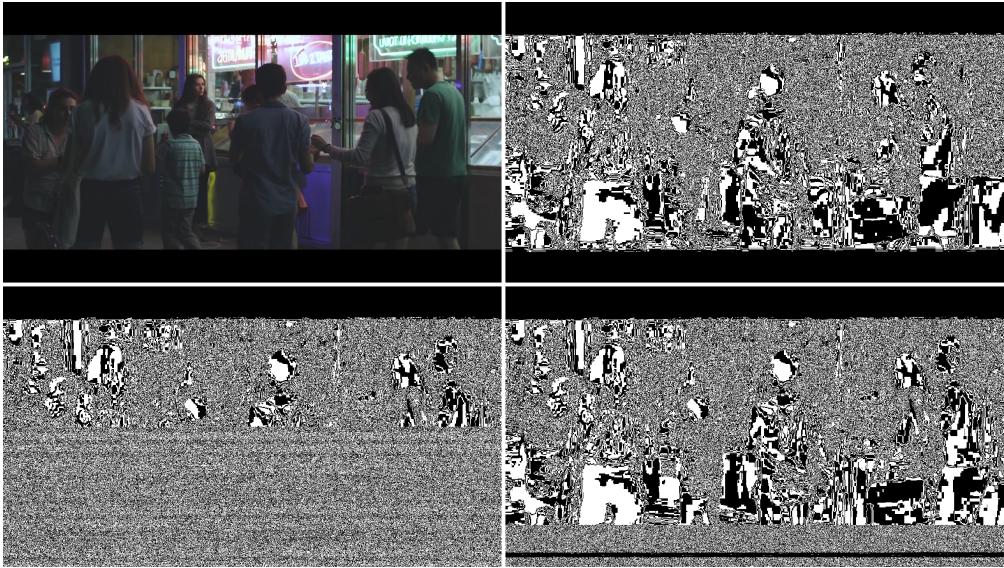


Figure 3.1: Illustration of the capacity parameter.

The **AVI Decoder** as described now provides all necessary functionality to allow the modification of pixel data to achieve the steganographic embedding of information within video frames.

3.3 Steganographic Algorithms

The concept of a generic steganographic embedding algorithm is initially developed and implemented as a **Steganographic Algorithm** interface. As defined within the background discussion of steganographic concepts, a steganographic system consist of a pair of functions providing embedding and extraction functionality. The interface will therefore need to declare two functions **embed** and **extract** which will embed and extract data respectively. The **Video Decoder** interface defined above allows access to video data via individual frame **Chunk** objects. It was therefore decided to define the **embed** and **extract** functions to operate on **Chunk** objects. The following interface definition was therefore used.

```
class SteganographicAlgorithm {
protected:
    string password;
    VideoDecoder *dec;
public:
    virtual void embed(Chunk *c, char *data, int dataBytes, int offset)=0;
    virtual void extract(Chunk *c, char *output, int dataBytes, int offset)=0;
    virtual void getAlgorithmCode(char out[4]) = 0;
};
```

Listing 3.4: Stego Algorithm interface (`steg/steganographic_algorithm.h:8`)

The `embed` function should be read as “embed `dataBytes` bytes from `data` into chunk `c` starting at an offset `offset` bytes into the frame”. Similarly, the `extract` function should be read as “extract `dataBytes` bytes from chunk `c` starting at an offset `offset` bytes into the frame and put them into `output`”. It was decided to place the error handling logic within the file system code. The steganographic algorithm implementations therefore do not contain any such logic and instead assume that pointer `data` does indeed point to at least `dataBytes` bytes and so on.

`getAlgorithmCode` simply returns a (unique) 4 character algorithm identifier which is used when users specify which algorithm they want to use.

3.3.1 LSB Sequential Embedding

Sequential LSB embedding is arguably the simplest steganographic algorithm. It works by replacing the LSBs of the cover object with the message bits producing the stego image. Algorithm 3.3 show pseudocode for the LSB embedding algorithm.

Algorithm 3.3 LSB embedding algorithm

```

1: for i ← 0 upto dataBytes - 1 do
2:   for j ← 7 downto 0 do
3:     if The jth significat bit of data[i] == 1 then
4:       frame[offset++] |= 1
5:     else
6:       frame[offset++] &= ~1
```

The matching extraction algorithm is shown in algorithm 3.4.

Algorithm 3.4 LSB extraction algorithm

```

1: for i ← 0 upto dataBytes - 1 do
2:   for j ← 7 downto 0 do
3:     output[i] |= ((frame[frameByte++] & 1) << j)
```

The actual implementation is relatively short and so is included below in listing 3.5.

```

virtual void embed(Chunk *c, char *data, int dataBytes, int offset) {
    char *frame = c->getFrameData();
    for (int i = 0; i < dataBytes; i++) {
        for (int j = 7; j >= 0; j--) {
            if (((1 << j) & data[i]) >> j) == 1) {
                frame[offset++] |= 1;
            } else {
                frame[offset++] &= ~1;
            }
        }
    }
};
```

Listing 3.5: LSB implementation (`steg/lsb_algorithm.cc:8`)

The operation of the sequential LSB embedding algorithm is illustrated in figure 3.2.

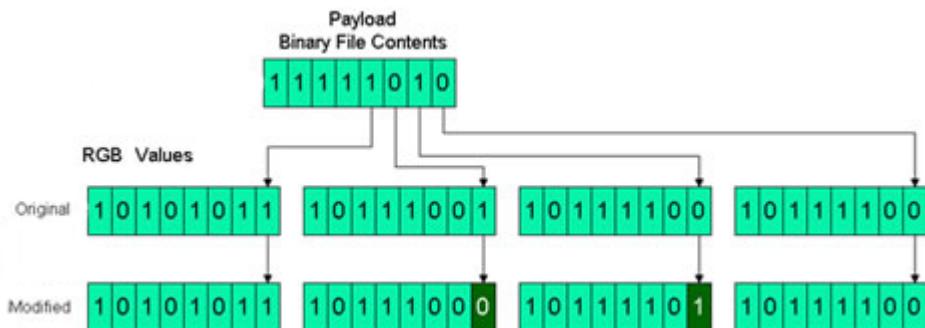


Figure 3.2: Illustration of the LSB Embedding algorithm.

Using the algorithm as shown above, 43.2 kB of data is embedded into a 1280×720 resolution video frame (a capacity setting of 100%). The result of this is shown in figure 3.3. The left image is the cover object and the right image is the stego object.



Figure 3.3: Effect of the LSB Embedding algorithm.

The visual impact on the video frame is very small and almost certainly not noticeable having been reproduced within this document at a smaller resolution. However, if we take a closer look at a specific portion of the video frame we can see some small discrepancies between the cover and stego objects - figure 3.4. I am unsure how well these images will be reproduced when printed, but the difference is definitely noticeable within the PDF. Note that without the original cover object for comparison, it would be very hard (if not impossible) to identify these details and deduce the presence of embedded data.

However, if the LSB plane of the frame is visually inspected, as in figure 3.5, one can immediately detect the presence of the hidden data - the LSB plane of the stego object looks “too random”. This is an example of a *visual steganalysis attack*.

Formalising the looking “too random” idea leads to another steganalysis attack known as the *Chi-Squared attack* developed by Westfeld and Pfitzmann. To implement the Chi-Squared attack, the concept of *Pair of Values* (PoVs) is first introduced.

One of the results of an embedding algorithm like sequential LSB embedding is the



Figure 3.4: Close up effect of the LSB Embedding algorithm.

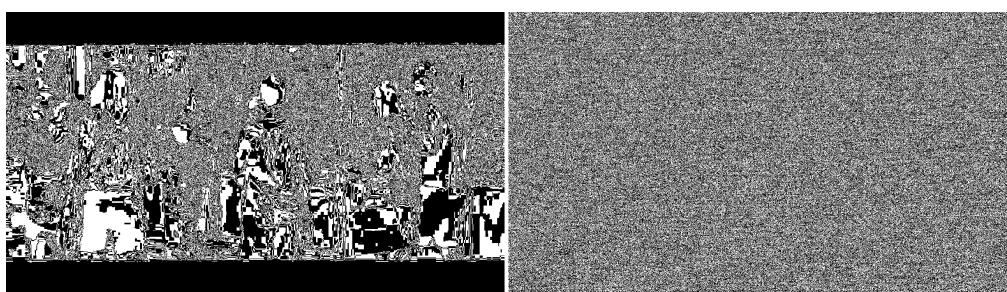


Figure 3.5: LSB plane of the cover and stego object.

creation of POVs, pixel values that embed into one another. For example, a pixel value of 100 in the cover image will either stay 100 or change to 101. Similarly, a pixel value of 101 will either stay 101 or change to 100. Thus (100, 101) is a POV.

Definition 3.1. PAIRS OF VALUES

A POV p is a member of the set \mathcal{P}_{ov} , defined as,

$$\mathcal{P}_{\text{ov}} \triangleq \{ (2k, 2k + 1) \mid 0 \leq k \leq 127 \}$$

Westfeld and Pfitzmann claim that the LSBs in images are not completely random, rather, the frequencies of each of the two pixel values in each POV tend to lie far from the mean of the POV. That is, it is unlikely for the frequency of pixel value $2k$ to be close to equal to the frequency of pixel value $2k + 1$. Furthermore, as information is embedded into the cover object, the frequencies of $2k$ and $2k + 1$ become (nearly) equal. The Chi-squared attack was designed to detect this and bases the probability of embedding on how close to equal POVs are in the image.

To implement the attack, the following steps are followed. First, $x_k = \text{frequency}(2k)$ and $y_k = \text{frequency}(2k + 1)$ are calculated $\forall k$, followed by the expected frequency $z_k = \frac{x_k + y_k}{2}$. n is defined to be the number of POVs, that is $|\mathcal{P}_{\text{ov}}|$. For uncompressed AVI using 24 bits per pixel, $n = 128$. The *minimum frequency condition* is now applied. This sets $x_k = y_k = z_k = 0$ and decrements n by one, if the condition $x_k + y_k \leq 4$ holds. The Chi-Squared statistic, with $n - 1$ degrees of freedom is then calculated:

$$\chi_{n-1}^2 = \sum_{k=0}^{127} \frac{(x_k - z_k)^2}{z_k}$$

The probability of embedding, p , is then calculated by evaluation of the following integral:

$$p = 1 - \frac{1}{2^{\frac{n-1}{2}} \Gamma(\frac{n-1}{2})} \int_0^{\lambda_{n-1}^2} e^{-\frac{u}{2}} u^{\frac{n-1}{2}-1} du$$

See appendix section ?? for an implementation of the Chi-squared attack in Python.

The Chi-squared attack produces very good results for the naïve sequential LSB embedding algorithm. Figure 3.6 show an example of it in use operating on a video frame with a capacity setting of 50%. It obvious that information has been embedded within the first half of the frame.

The Chi-Squared attack motivates the development of a *permuted* LSB embedding algorithm. The attack works so well because regardless of the capacity setting, the data is sequentially embedded from the top of the video frame. It would be better to distribute the data uniformly throughout the video frame. A permuted LSB embedding algorithm achieves this and resists the Chi-Squared attack when a small enough capacity setting is used²⁰.

²⁰Of course, a permuted embedding using a capacity of 100% is comparatively no different to sequential embedding!

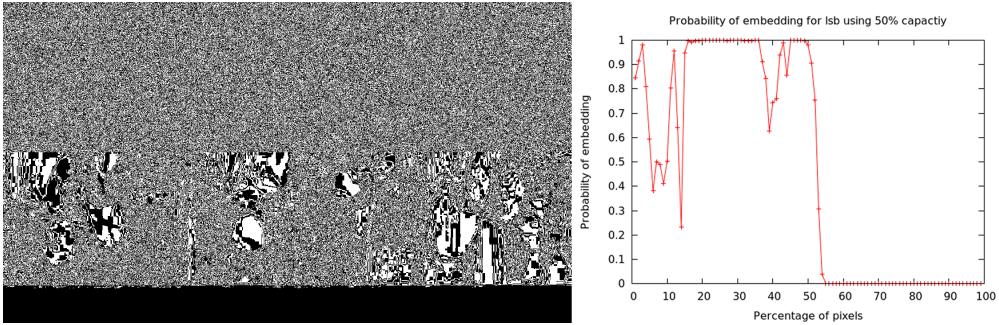


Figure 3.6: Results of the Chi-Squared attack.

3.3.2 Permuted LSB Embedding

One can produce deceptively simple pseudocode for the permuted LSB embedding algorithm as shown in algorithm 3.5.

Algorithm 3.5 Permuted LSB embedding algorithm

```

1: path ← a pseudorandom permutation of the cover object
2: path.seekToOffset(offset-1)
3: for i ← 0 upto dataBytes - 1 do
4:   for j ← 7 downto 0 do
5:     if The jth significat bit of data[i] == 1 then
6:       Set LSB(frame[path.next()]) to 1
7:     else
8:       Set LSB(frame[path.next()]) to 0

```

Lines 1 and 2 hide a large amount of complexity involved in actually implementing the permuted LSB algorithm.

Assuming the cover object consists of n bytes, a pseudorandom permutation of the cover object can be thought of as a pseudorandom permutation of the numbers 0 - n . The obvious way to produce a pseudorandom permutation of a list of numbers is to shuffle an array containing them. This approach has the drawback that you need to hold the numbers in memory (11 MB for a single 720p HD video frame). Instead, a different approach using a *Linear Congruential Generator* (LCG) was taken.

Definition 3.2. LINEAR CONGRUENTIAL GENERATOR

A *Linear Congruential Generator* is defined by the recurrence relation:

$$X_{n+1} = (aX_n + c) \bmod m$$

where X is the sequence of pseudorandom values, and

$$\begin{aligned} m, & 0 < m - \text{the modulus} \\ a, & 0 < a < m - \text{the multiplier} \end{aligned}$$

$$c, \quad 0 \leq c < m - \text{the increment} \\ X_0, \quad 0 \leq X_0 < m - \text{the seed}$$

are integer constants that specify the generator.

The Hull-Dobell Theorem states that a LCG will have a full period if and only if the following 3 requirements are satisfied:

1. $\gcd(c, m) = 1$ (c and m are relatively prime),
2. $a - 1$ is divisible by all prime factors of m ,
3. $a - 1$ is a multiple of 4 if m is a multiple of 4.

Therefore, if the above requirements can be satisfied, a LCG can be used as a pseudorandom permutation for the cover object with m equal to n (the size of the cover object). Note that forcing m to be a power of 2, simplifies the above requirements to:

1. $\gcd(c, m) = 1$,
2. $a - 1$ is odd,
3. $a - 1$ is a multiple of 4 if m is a multiple of 4.

As such, m is set equal to the next power of 2 larger or equal to n . Any produced values larger or equal to n are just thrown away, therefore producing a full period of size n as required.

The pseudorandom permutation should be dependant on a user provided password. Therefore, the values of a and c are determined using a key derived from the users password. The popular key derivation function PBKDF2 is used along with the **Whirlpool** hash function to generate a 128 byte key and the first 4 bytes are taken as an unsigned integer and used to derive c and a .

This successfully describes an implementation for line 1 of algorithm 3.5. Line 2 unfortunately causes some problems for the LCG due to precisely the reason that it was chosen to be used.

Since the full array of the pseudorandom permutation is not being stored, the LCG does not facilitate direct random access. For example, if the 50th element of the permutation is requested, the 49 elements before it must first be calculated starting from X_0 . This had a huge impact on the performance of the algorithm to the point of the file system becoming unusable. Therefore, it was decided to pre-compute a hashmap mapping from frame offsets to the corresponding pseudorandom permutation element, therefore eliminating the need to calculate all elements prior to the requested one. This however now means the entire permutation is being stored in memory - exactly why the original method was rejected! It was decided to remain using the LCG method with only a single global permutation used across all video frames.

The result of the permuted embedding algorithm is illustrated in figure 3.7. The left image has data embedded using the sequential embedding algorithm whereas the right

image has the same data embedded using the permuted embedding algorithm. Both images were using a capacity setting of 15%.

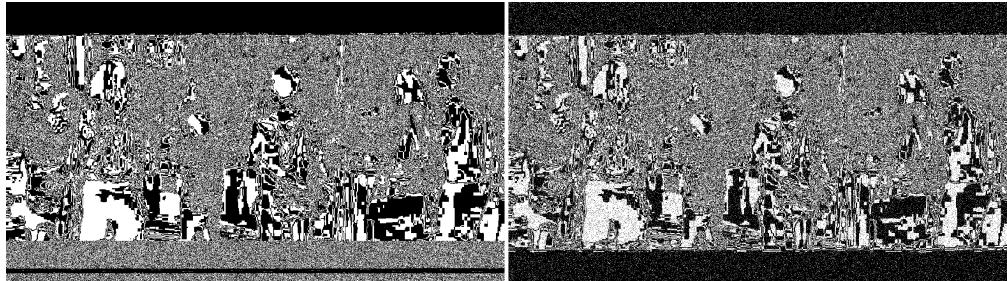


Figure 3.7: Illustration of the permuted LSB algorithm.

Figure ?? show the result of applying the Chi-Squared attack to a video frame with data embedded using the sequential LSB algorithm (left) and the permuted LSB algorithm (right), both using a capacity setting of 15%.

3.3.3 Permutated and Encrypted LSB Embedding

To further strengthen the security block.

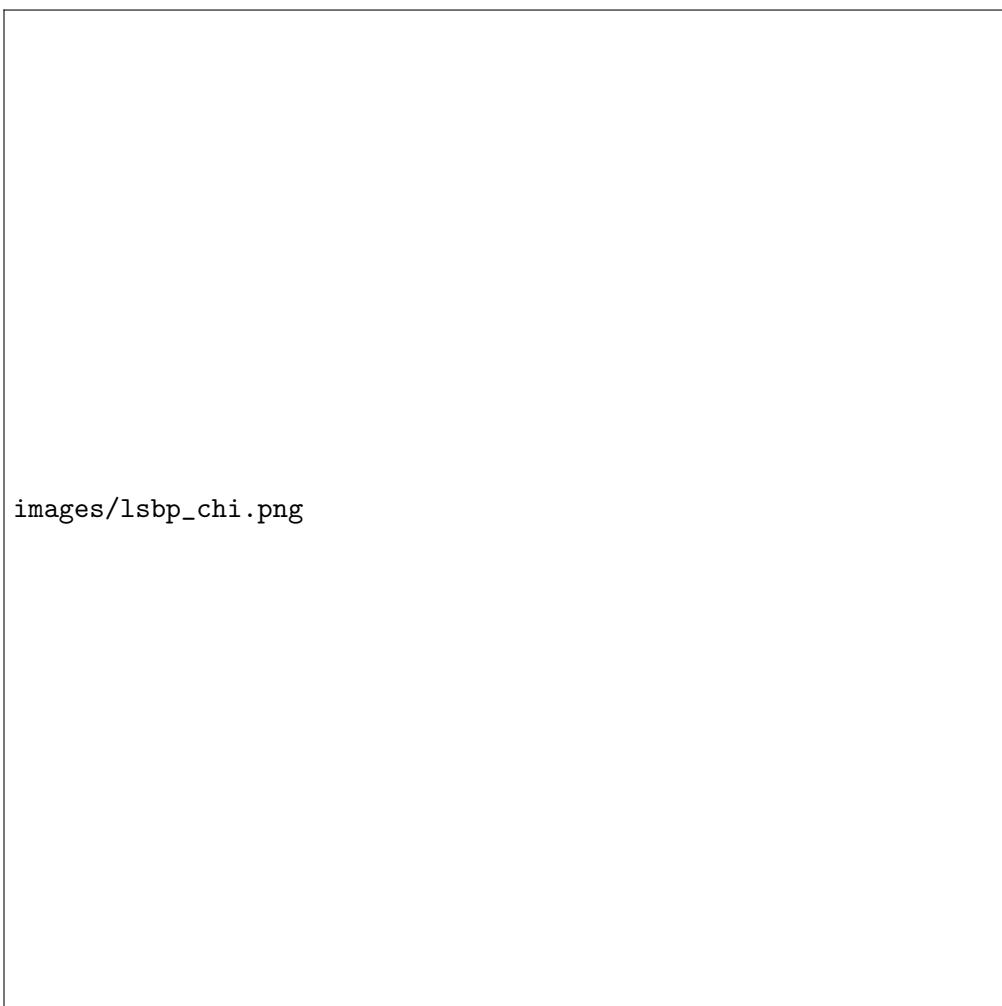


Figure 3.8: Chi-Squared attack on the permuted LSB algorithm.

3.4 File system

3.5 Extensions

3.5.1 Supporting multiple video formats

3.5.2 File system directory structures

3.5.3 Plausible deniability

3.5.4 Hiding data within audio streams

3.6 Testing

4 || Evaluation

4.1 Satisfaction of Requirements

4.2 Correctness

4.3 Security

4.4 Performance

5 || Conclusions

5.1 Future Project Directions

References

[1] *Steganography in Digital Media*. Jessica Fridrich, 2010.