

Scott Williams

STEGANOGRAPHIC FILE SYSTEMS WITHIN VIDEO FILES

COMPUTER SCIENCE PART II PROJECT DISSERTATION

Christ's College
University of Cambridge

May 1, 2015

Proforma

NAME:	Scott Williams
COLLEGE:	Christ's
PROJECT TITLE:	Steganographic file systems within video files
EXAMINATION:	Part II of the Computer Science Tripos
YEAR:	2015
WORD COUNT:	11,957
PROJECT ORIGINATOR:	Scott Williams
PROJECT SUPERVISOR:	Daniel Thomas

Original Aims of the Project

To investigate appropriate steganographic embedding methods for video and to develop a practical steganographic software package to enable the embedding of arbitrary data within video files via a file system interface. Raw AVI video files should be supported and a variety of steganographic embedding algorithms should be available. Basic file system commands should work within the presented logical volume and embedding should occur with no perceivable impact on video quality. If time permits, multiple video formats should be supported along with encryption and plausible deniability functionality.

Work Completed

A complete software package has been developed enabling the embedding of arbitrary files within multiple video formats via a file system interface. A native uncompressed AVI decoder has been implemented along with a generic method for dealing with other video formats. A total of 6 steganographic embedding algorithms are supported, along with encryption and plausible deniability functionality. Common file system operations (including directories) work as expected within the mounted volume and the embedding process can operate without any perceivable impact on video quality as verified by a user study. The performance of the file system is adequate for general use achieving read and write speeds on par with USB 2.0 flash drives.

Special Difficulties

None.

Declaration of Originality

I, Scott Williams of Christ's College, being a candidate for Part II of the Computer Science Tripos, hereby declare that this dissertation and the work described in it are my own work, unaided except as may be specified below, and that the dissertation does not contain material that has already been used to any substantial extent for a comparable purpose.

I give permission for my dissertation to be made available in the archive area of the Laboratory's website.

Signed:

Date:

Contents

1	Introduction	1
1.1	Motivation	1
2	Preparation	2
2.1	Background	2
2.1.1	Steganographic Concepts	2
2.1.2	Steganalysis	4
2.1.3	The AVI file format	5
2.1.4	JPEG compression	6
2.1.5	FFmpeg	6
2.1.6	Developing a file system	7
2.2	Existing tools	8
2.2.1	StegoStick	8
2.2.2	StegoMagic	9
2.2.3	TCSteg	9
2.2.4	StegoVideo	9
2.2.5	OpenPuff	9
2.2.6	Steganosaurus	10
2.3	Requirements Analysis	10
2.3.1	Core Requirements	10
2.3.2	Possible Extensions	11
2.4	Choice of Languages and Tools	12
3	Implementation	14
3.1	Introduction	14
3.2	AVI Decoder	14
3.3	Steganographic Algorithms	17
3.3.1	LSB Sequential Embedding	18
3.3.2	Permuted LSB Embedding	21
3.3.3	Encrypting the embedded data	24
3.4	The file system	25
3.4.1	Developing the header	26
3.4.2	Writing to the file system	27
3.4.3	Reading from the file system	28
3.4.4	Listing files in the file system	28
3.5	Command line application	29
3.6	Extension Tasks	29
3.6.1	Supporting multiple video formats	29
3.6.2	File system directory structures	33
3.6.3	Plausible deniability	34

3.6.4	Hiding data within audio	35
3.6.5	Evaluation of embedding impact on video quality	35
3.7	Testing	36
4	Evaluation	38
4.1	Satisfaction of Requirements	38
4.1.1	Embedding data within video files	38
4.1.2	Providing a file system interface	38
4.1.3	Supporting uncompressed AVI	38
4.1.4	File system performance	38
4.1.5	Supporting multiple video formats	39
4.1.6	File system directory operations	39
4.1.7	Embedding within audio	39
4.1.8	Plausible deniability	39
4.1.9	Evaluation of the visual impact of embedding	39
4.2	Security	39
4.3	Performance	41
4.3.1	Performance of the file system	42
5	Conclusions	44
5.1	Lessons Learnt	44
5.2	Future Project Directions	44
	References	45
A	Details of the AVI file format	49
A.1	Detailed AVI form	49
A.2	The AVI and Bitmapinfo headers	50
B	Stegasis example use	51
B.1	Stegasis usage information	51
B.2	Example use case	52
C	Detailed Code samples	54
C.1	Chi-Squared attack	54
C.2	Reading and writing to the file system	56
C.3	Implementing F5	57
D	Testing	61
D.1	Unit Testing	61
D.2	Integration Testing	61
E	User study results	63
F	Original Project Proposal	64

1 || Introduction

Steganography is the art of hiding information in apparently inconspicuous objects. Whereas cryptography seeks to protect the content of information, steganography attempts to conceal the fact that the information exists^[1, pp. xv - xvi]. This allows steganographic methods to be utilised in countries where encryption is illegal and identification of encrypted data can be grounds for imprisonment. Within the UK for example, keys for identified encrypted data can be forced to be disclosed^[2], rendering standard cryptographic methods alone unfavourable.

In this project I design and implement a practical steganographic software application - **Stegasis** - which enables users to embed arbitrary files within video via a file system interface. **Stegasis** can operate with no perceivable impact on video quality and can provide large embedding capacities. Multiple video formats are supported along with several steganographic embedding algorithms. Standard encryption algorithms can be used to further protect embedded data and plausible deniability functionality keeps sensitive information safe even when the presence of embedded data has been confirmed.

Steganographic methods operating on video have had comparatively little attention compared to other medium such as images and audio^[1, p. xvii]. As such, there are few programs currently available which allow data to be steganographically hidden within video^[3]. To the best of my knowledge, **Stegasis** is the first application to enable the embedding of arbitrary files within videos via a file system interface.

1.1 Motivation

Digital media is ubiquitous on the Internet and high definition video content is now common place on video sharing and social networking websites. Video files of multiple gigabytes in size can reside on users devices without arousing suspicion, providing an ideal hiding place for large collections of sensitive files. Few programs are capitalising on this fact, and those that are, allow the user to embed only a single chosen file into a small range of video formats with very low embedding capacities. As with **TrueCrypt**¹, I believe that a practical system for protecting sensitive files should present the user with a mounted logical volume allowing the use of standard file system operations to create, access and organise embedded data. Furthermore, there exist many commonly used video formats along with many more currently in development. As such, a steganographic program operating on a small number of video formats not only greatly restricts usability, it will require constant development as new video formats inevitably become more popular. Instead, a generic solution applicable to a variety of video formats is preferred.

The recent global surveillance disclosures show the extent to which government authorities monitor online communications^[5]. These, together with current UK laws mean it is no longer the case that simply encrypting data is enough to keep the owner safe.

¹A successful widely used cryptographic program providing on-the-fly and full disc encryption^[4]. Unfortunately, **TrueCrypt** is no longer being maintained.

2 || Preparation

2.1 Background

The project aims to combine 3 main topics; steganography, video formats and file systems. Understanding of these will be required to develop a successful final product.

The most important property of any steganographic system is undetectability, that is, it should be impossible to differentiate between ordinary and steganographically modified objects. This requirement is famously formulated within Simmons' prisoners' problem^[6]:

Alice and Bob are imprisoned in separate cells and wish to formulate an escape plan. They are allowed to communicate, but all messages must pass through a warden Eve. If Eve suspects the prisoners of secretly discussing their escape plan, the communication channel will be severed and Alice and Bob thrown into solitary confinement. The prisoners attempt to utilise steganography to exchange details of their plan undetected. The steganographic system is considered broken if Eve is able to detect the presence of hidden messages within the prisoners exchanges. It is assumed that Eve has a complete knowledge of the steganographic algorithm being used, with the exception of the stego key, which Alice and Bob have agreed upon beforehand. This is in parallel with Kerckhoff's principle used within cryptography^[7]. The warden can be considered to be: *passive*, *active* and *malicious*. A passive warden does not modify the exchanged messages in any way, whereas an active warden may modify the messages whilst maintaining their original meaning. For example replacing words with synonyms or reordering sentences. If images are being used as a transport medium then an active warden may recompress or crop the images. A malicious warden attempts to break the steganographic system and impersonate the prisoners in an attempt to obtain information.

This project is concerned only with the case of the passive warden. As such, any modification of the video files once **Stegasis** has embedded data within them, will most likely render the embedded file system corrupt.

2.1.1 Steganographic Concepts

A *steganographic system* will form a core part of the final application, allowing data to be hidden within video. A steganographic system depends on the following components:

- A *Cover object* is the original object that the message will be embedded within. A cover object consists of a number of *elements*, for example pixels.
- A *Message* is an arbitrary length sequence of symbols. For this project we consider messages of the form $\mathcal{M} \in \{0, 1\}^{8 \cdot n}$ for some n - a sequence of bytes.
- A *Stego key* is a secret key used within the embedding process.
- A *Stego object* is the result of embedding a message inside a cover object.

Definition 2.1. STEGANOGRAPHIC SYSTEM

Let \mathcal{C} be the set of all cover objects. For a given $\mathbf{c} \in \mathcal{C}$, let $\mathcal{K}_{\mathbf{c}}$ denote the set of all stego keys for \mathbf{c} , and the set $\mathcal{M}_{\mathbf{c}}$ denote all messages that can be communicated in \mathbf{c} . A steganographic system² is then formally defined as a pair of embedding and extracting functions *Emb* and *Ext*,

$$\begin{aligned} \text{Emb} : \mathcal{C} \times \mathcal{K} \times \mathcal{M} &\rightarrow \mathcal{C} \\ \text{Ext} : \mathcal{C} \times \mathcal{K} &\rightarrow \mathcal{M} \end{aligned}$$

satisfying,

$$\forall \mathbf{c} \in \mathcal{C}, \mathbf{k} \in \mathcal{K}_{\mathbf{c}}, \mathbf{m} \in \mathcal{M}_{\mathbf{c}}. \text{Ext}(\text{Emb}(\mathbf{c}, \mathbf{k}, \mathbf{m}), \mathbf{k}) = \mathbf{m}$$

Definition 2.2. EMBEDDING CAPACITY

The Embedding Capacity (payload) $\mathcal{P}_{\mathbf{c}}$ for a given cover object $\mathbf{c} \in \mathcal{C}$ is defined in bits as,

$$\mathcal{P}_{\mathbf{c}} = \log_2 |\mathcal{M}_{\mathbf{c}}|$$

The relative embedding capacity $\mathcal{R}_{\mathbf{c}}$ for a given cover object $\mathbf{c} \in \mathcal{C}$ is defined as,

$$\mathcal{R}_{\mathbf{c}} = \frac{\log_2 |\mathcal{M}_{\mathbf{c}}|}{n}$$

where n is the number of elements in \mathbf{c} .

For example, consider \mathcal{C} to be the set of all 512×512 greyscale images, embedding one bit per pixel gives $\mathcal{M} = \{0, 1\}^{512 \times 512}$ and $\forall \mathbf{c} \in \mathcal{C}. |\mathcal{M}(\mathbf{c})| = 2^{512 \times 512}$. The embedding capacity $\forall \mathbf{c} \in \mathcal{C}$ is then $512 \times 512 \approx 33\text{kB}$ as expected. In this case, n is equal to the number of pixels in \mathbf{c} and therefore the relative embedding capacity is equal to 1 bpp (bits per pixel), again as expected.

Using the definitions above, we can define a simple expression for the embedding capacity of a video file.

Definition 2.3. EMBEDDING CAPACITY FOR VIDEO

With \mathcal{C} as the set of all video files, the embedding capacity $\mathcal{V}_{\mathbf{c}}$ for a given video $\mathbf{c} \in \mathcal{C}$ can be expressed as,

$$\mathcal{V}_{\mathbf{c}} = \sum_{f \in \text{frames}(\mathbf{c})} \mathcal{P}_f$$

Note that for certain embedding algorithms, the embedding capacity can depend on both the input data and the cover object³. However, in some cases the following expression is also valid,

$$\mathcal{V}_{\mathbf{c}} = |\text{frames}(\mathbf{c})| \cdot \mathcal{P}_{f_0}$$

²This is specifically steganography by cover modification.

³Many algorithms operating on JPEG images for example will not embed within zero valued DCT coefficients.

Definition 2.4. STEGANOGRAPHIC CAPACITY

The concept of Steganographic Capacity is loosely defined as the maximum number of bits that can be embedded within a given cover object without introducing statistically detectable artifacts.

For completeness, the least significant bit (LSB) of a given number is defined as follows,

$$\text{LSB}(x) = x \bmod 2$$

It will be useful to visually inspect the effect of steganographic embedding algorithms operating on the LSBs of pixels. The *LSB Plane* of an image is therefore defined.

Definition 2.5. LSB PLANE

The Least Significant Bit Plane of a given image \mathbf{c} and a specified colour channel q is defined as the 1 bit image $\text{LSBP}(\mathbf{c}, q)$ which has resolution equal to that of image \mathbf{c} and with pixel values $\text{LSBP}(\mathbf{c}, q)(x, y)$ given by,

$$\text{LSBP}(\mathbf{c}, q)(x, y) = \text{LSB}_q(\mathbf{c}(x, y))$$

2.1.2 Steganalysis

Steganalysis is the study of detecting messages embedded using steganographic techniques; this is analogous to cryptanalysis applied to cryptography^[8]. A steganalysis attack is considered successful (that is, the steganography has been broken) if it is possible to correctly distinguish between cover and stego objects with probability better than random guessing. Note that it is not necessary to be able to read the contents of the secret message to break a steganographic system.

A trivial example of steganalysis arises when the steganalyst has access to the original cover object used within the embedding procedure. By computing the difference between the stego and cover objects, the steganalyst can immediately detect the presence of a hidden message. This attack identifies a number of important points to consider when developing a practical steganographic system. Firstly, embedding within popular media content should be discouraged, as the cover object will likely be widely available. Secondly, if a user is embedding within original content, for example a video recorded by them, any copies of the original file should be securely erased after embedding.

Steganalysis methods can be split into two main categories, *Targeted Steganalysis* and *Blind Steganalysis*. Targeted Steganalysis occurs when the steganalyst has access to the details of the steganographic algorithm used for embedding. The steganalyst can accordingly target their activity to the specific stegosystem. On the other hand, if the steganalyst has no knowledge of the utilised steganographic algorithm, Blind Steganalysis techniques must be applied. In this project, Targeted Steganalysis attacks are developed for several of the proposed embedding algorithms.

2.1.3 The AVI file format

As specified within the project proposal, this project initially looks at raw uncompressed AVI files. Furthermore, only AVI version 1.0⁴ files are investigated and therefore supported natively⁵ by **Stegasis**. Unfortunately, uncompressed AVI is today, a very uncommon video format^[9]. This is likely due to its relatively huge file sizes when compared to a modern compressed format such as H.264. For example, one minute of 720p HD footage encoded as uncompressed AVI is roughly 4.2 GB.

The AVI file format is a Resource Interchange File Format (RIFF) file specification developed by Microsoft and originally introduced in November 1992^[10]. The data within RIFF files is divided into chunks and lists, each of which is identified by a FourCC tag. An AVI file takes the form of a single chunk in a RIFF formatted file, which is then subdivided into two mandatory lists, the `hdr1` and `movi` and one optional chunk, the `idx1`. The second sub-list contains the actual audio/video data and will be where steganographic embedding will occur. See the Appendix Section A for detailed definitions of the data structures used.

An AVI file consists of a number of data streams (usually 2, one for audio and one for video) interleaved within the `movi` list. Each stream has a corresponding AVI stream header and format chunk within the `hdr1` list. These data structures contain information about the stream including the codec and compression used (if any). Specifically, the `fccHandler` field contains a FourCC tag that identifies a specific data handler. For raw uncompressed video this will equal 'DIB ' (Device Independent Bitmap). Any user provided AVI files with a `fccHandler` not equal to 'DIB ', (compressed video) will at this point be rejected and an error message presented to the user.

The `movi` list contains the raw video and audio data within sequential RIFF chunks. Each chunk for the DIB video stream contains one frames worth of pixel data, with each pixel represented by a 3 byte BGR (Blue Green Red) triple - a total of 24 bits per pixel. The first 3 byte triple corresponds to the lower left pixel of the final image⁶.

If we use an embedding algorithm which embeds 3 bits per pixel (1 bit per colour channel per pixel), we can derive a simple expression for the embedding capacity of an uncompressed AVI video c , in terms of the height h and width w in pixels and the total number of frames t :

$$\mathcal{V}_c = 3 \cdot w \cdot h \cdot t$$

These values are all available within the `AVIMAINHEADER` structure allowing the user to be informed of the video's total embedding capacity upon formatting.

⁴Not including the Open-DML extension (version 1.02).

⁵All other video formats (including compressed AVI) are supported via the use of **FFmpeg**, as described in Section 3.6.1.

⁶This can be inverted via a flag within the `BITMAPINFOHEADER`.

2.1.4 JPEG compression

The JPEG file format will prove useful when developing a universal steganographic technique operating across many video formats (Section 3.6.1). Steganography within the JPEG format has had a comparatively large amount of attention from the research community^[11]. As such, there exists a fair number of well documented steganographic embedding algorithms for JPEG^{[12] [40] [13]}.

The JPEG compression process consists of 5 main procedures:

1. Transform the image into an optimal color space.
2. Downsample chrominance components by averaging groups of pixels together.
3. Apply a *Discrete Cosine Transform* (DCT) to blocks of pixels.
4. Quantise each block of DCT coefficients using a quantisation table.
5. Encode the resulting coefficients using a Huffman variable word-length algorithm.

Note that step 4 is an example of lossy compression, whereas step 5 is lossless. Therefore most steganographic algorithms will operate on the quantised DCT coefficients (between steps 4 and 5) to avoid embedded data being lost due to quantisation.

Conveniently, the Independent JPEG Group provide the `libjpeg` C library^[14] which will abstract the complexities of the JPEG format and allow direct access to the quantised DCT coefficients prior to step 5 being executed.

JPEG DCT coefficients are arranged into several *components* containing *rows* which contain a number of *blocks*. Each block contains 64 coefficients ranging from -2048 to 2047 . A JPEG image will usually have 3 components corresponding to the *Luminance and Chrominance* colour model (YCbCr) with the first component being luma and the remaining two being the blue-difference and red-difference chroma components. Since human perception is more sensitive to changes in luminance compared to colour^[15], steganographic embedding will usually not occur within the luminance component.

It is worth noting that the JPEG decompression and compression processes are computationally expensive. This is especially important when dealing with video since a 3 minute video, for example, consists of around 4,500 frames (which can be considered as individual JPEGs). Since performance of the virtual file system is important, design decisions will need to be made to avoid any unnecessary compression and decompression operations. Additionally, although JPEG files are small on disk, they are not once decompressed. It will not be possible to hold all decompressed JPEG frames of a video in memory.

2.1.5 FFmpeg

FFmpeg is an open source, multimedia framework^[16]. It is a “complete, cross-platform solution to record, convert and stream audio and video”. In particular, it contains codecs for nearly every video format available today^[17].

The pitfalls of uncompressed AVI mean that **Stegasis** would greatly benefit from operating on multiple video formats. I could continue to investigate more video formats

and develop parsers for these as part of the project. However, this will become a very time consuming endeavor most likely resulting in very brittle, untested parsers. Instead, it would be wise to leverage the **FFmpeg** framework for this functionality.

A novel solution to this problem is posed in Section 3.6.1 and makes use of **FFmpeg** to convert input files to the *Motion JPEG* file format.

2.1.6 Developing a file system

A file system can either operate within *kernel* or *user space*. I decided within the project proposal to develop the file system component for **Stegasis** in user space using the **FUSE** (Filesystem in Userspace) library^[18]. This was primarily because developing a kernel module is complex and hard to test - a segmentation fault occurring within kernel space code will bring down the entire machine. A kernel module also requires a large amount of boiler plate code and I would prefer to spend time on the steganographic portion of this project rather than getting bogged down with the complexities of a kernel file system implementation. In contrast, **FUSE** ships with an example “hello world” file system which is less than 100 lines of C code. Also, developing the file system in user space will cause the final application to be a lot more portable and easier for users to install - a kernel space file system would require super user permission to load the related kernel module.

There are however disadvantages to using a file system in user space, performance being one of them^[19]. This is due to the **FUSE** kernel module having to act as a proxy between the system call and the user space code.

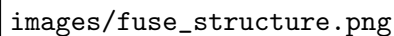


Figure 2.1: Path of a file system call in FUSE, taken from `fuse.sourceforge.net`.

Figure 2.1 shows the path of a file system call in the provided hello world example file system. We can see the FUSE kernel module acting as a proxy between the VFS system call and the `example/hello` user space code. A kernel space file system would not need to re-enter user space to complete the system call, hence giving better performance.

The FUSE library provides a number of function definitions which the user space code implements. These functions are then called when the corresponding file system operation occurs.

```
int read(const char *path, char *buf, size_t size, off_t offset, struct
        fuse_file_info *fi);
```

Listing 2.1: FUSE read operation.

The `read` function is called when a file system read occurs. It requests that `size` bytes of the file `path` starting at offset `offset` should be written to the buffer `buf`.

2.2 Existing tools

The relatively little work on steganography within video was reflected in my search for steganographic programs operating on video files. This section contains an exhaustive list of all the video steganography tools I could find freely⁷ available on the Internet. A total of 6 tools claimed to provide steganographic embedding functionality within video files. Of these 6, only 3 actually attempt to embed within the video data itself. None of the identified programs allow the user to embed more than one file⁸ and none of them provide any sort of file system interface.

2.2.1 StegoStick

StegoStick^[20] claims to allow users to “hide any file into any file”. This statement suggests that the program is simply appending the requested file to the end of the cover object. This suspicion is partly true; based on the file extension, **StegoStick** splits cover objects into 3 categories: images, media and other. The other category does indeed just append the file to the cover object, whereas the image and media category do attempt to employ steganographic embedding methods. The images category applies to **BMP** files and uses LSB embedding. The media category applies to **WAV**, **AVI** and **MPG** files and assumes each format has a “header” of 99 bytes⁹. Although this seems to be true for the **WAV** format, this is not the case for **AVI** nor **MPG** files. **StegoStick** will then use blind LSB embedding within the remaining data. As such, my attempts to use **StegoStick** to embed within **AVI** files rendered the resulting video unplayable.

⁷A further 2 programs exist claiming to embed within video, however these are not freely available to download (Info Stego, Hiderman).

⁸Admittedly you could embed a compressed archive using these tools to effectively allow a directory structure to be embedded.

⁹Listed in the source as “44 byte header + 54 bytes of extension space”.

2.2.2 StegoMagic

StegoMagic^[21] claims to “work on all types of files and all size of data” which again sounds as though it’s appending the file to the end of the cover object. This is indeed the case, embedding an image within a video and inspecting the modified file shows that data has just been appended to the end of the video, albeit encrypted. **StegoMagic** does not specify the encryption algorithm used and the source code is not available to view. Furthermore, the user cannot specify an encryption key to use. Instead, **StegoMagic** generates a 5 digit number during the embedding process and presents this to the user.

2.2.3 TCSteg

TCSteg^[22] is a Python script accompanying a blog post written by Martin Fiedler discussing hiding **TrueCrypt** volumes within MP4 files. The method described embeds the **TrueCrypt** volume within the MP4 atom **mdat** and modifies the chunk offset table within the **moov** atom so that any application playing the video will ignore the embedded data. A nice property of **TCSteg** is that the resulting video file can be directly mounted by **TrueCrypt** since it ignores the MP4 header data prior to the embedded volume.

The above programs all resort to embedding within video files by either appending the embedded data to the end of the video, or inserting the embedded data at some point within the video file. I do not consider this approach to embedding data secure, and it should be a trivial task for any steganalyst to detect the presence of embedded data within the stego objects using a simple hex editor. Therefore, the above stegosystems should not be used for hiding sensitive data.

2.2.4 StegoVideo

StegoVideo^[23] is a Virtual Dub filter¹⁰ which allows users to embed a file within AVI files (supporting multiple compression codecs). I am unsure of the exact steganographic embedding algorithm used since the program is closed source, but the website does mention that **StegoVideo** makes use of error correction codes to allow embedded data to be recovered even after the resulting video has been compressed - although this is understandably dependant on the compression amount. **StegoVideo** attempts to protect the embedded data via the use of a 5 digit number, although as with **StegoMagic**, this is not provided by the user and is instead generated and presented to the user.

2.2.5 OpenPuff

OpenPuff^[24] is a steganographic tool supporting a wide range of formats, including 3GP, MP4, MPG and VOB. It allows users to embed a file within a collection of carrier objects and uses 3 user provided passwords to encrypt, scramble and whiten (mixing with noise) the provided file. Plausible deniability is also provided via the option to add decoy content.

¹⁰Which is also available in a stand alone executable form.

`OpenPuff` successfully embedded and retrieved a text file within a sample MP4 video and I could notice no perceivable impact on video quality. Performance was also good due to multithreading support. However, the embedding capacity is very limited. I was only able to achieve embedding capacities of around 0.0043%¹¹ even at the maximum capacity setting. This makes `OpenPuff` impractical for hiding large files - requiring hundreds of carrier files to embed an average 2 MB JPEG image.

2.2.6 Steganosaurus

`Steganosaurus`^[25] is a cross platform steganographic program developed by James Ridgway. It allows users to embed a file within h.264 video files via the modification of motion vectors. Two embedding algorithm variants are provided and the input file is encrypted using AES with a user provided passphrase. I unfortunately could not get `Steganosaurus` to run on my computer (using Linux or Windows) and therefore could not test its operation.

The above 3 programs offer much more promising steganographic security and some of them also support multiple video formats. However, all feature the same limitation of only allowing the user to embed one chosen file and the offered embedding capacities are far from practical for use with large files.

This project aims to remedy these issues by allowing the user to embed an arbitrary number of files within a video (via a file system interface) and by providing high capacity steganographic embedding algorithms (for example 15% of the video size)¹².

2.3 Requirements Analysis

After reviewing the necessary background material and investigating current available solutions to the problem of steganography within video, I produced the following collection of requirements. For the project to be a success, all of the core requirements should be fulfilled.

2.3.1 Core Requirements

`Stegasis` should:

1. Allow users to embed data within video files:
 - a) Several steganographic embedding algorithms should be available.
 - b) Each embedding algorithm, \mathcal{A} , should satisfy correctness. That is,
$$\forall \mathbf{c}, \mathbf{k}, \mathbf{m}. \text{Ext}_{\mathcal{A}}(\text{Emb}_{\mathcal{A}}(\mathbf{c}, \mathbf{k}, \mathbf{m}), \mathbf{k}) = \mathbf{m}.$$
 - c) Embedding should occur with no perceivable impact on video quality.

¹¹2,600 bytes within a 60 MB video.

¹²This is very much a trade off - larger embedding capacities will come at the sacrifice of steganographic security. However, this decision is presented to the user with sensible defaults.

- d) Steganalysis tools should be developed to test the security of the proposed embedding algorithms.
 - e) An optional provided passphrase should encrypt data prior to embedding.
 - f) A capacity flag should allow users to specify the percentage of each video frame to embed within.
2. Provide a file system interface:
- a) The presented logical volume should reside at a user provided mount point.
 - b) Data written to the file system should be embedded on the fly within the chosen video file.
 - c) Data accessed from the file system should be extracted on the fly from within the video.
 - d) Standard file system operations such as creating, deleting and moving files should work as expected, and standard Unix tools such as `cp`, `mv` and `rm` should also work as expected on files within the volume.
3. Support raw uncompressed AVI video:
- a) Uncompressed AVIs should be natively decoded allowing access to individual pixel data.
4. Provide adequate file system performance:
- a) Full HD video content should be playable directly from within the presented file system (bitrates of full HD video are roughly 8 - 12 Mb/s^[26]).
 - b) Ideally, the file system should provide read and write speeds comparable to those provided by USB 2.0 devices¹³(roughly 20 MB/s^[27]).

2.3.2 Possible Extensions

If time constraints allow, the following extension tasks shall also be completed.

`Stegasis` should:

1. Support embedding within multiple video formats.
2. Allow directory operations within the file system:
 - a) Creating directories using the `mkdir` command should work as expected, as should using the `mv` and `rm` commands.
 - b) Organising files within directories should work as expected.
3. Embed also within audio data:

¹³Although the USB 2.0 standard supports speeds of up to 480 Mb/s, devices never reach this theoretical limit.

- a) Data should also be embedded within the (possible) audio stream of the video, therefore increasing the embedding capacity.
- 4. Provide plausible deniability:
 - a) A second file system should be (optionally) embedded within the video, mountable with a second passphrase.
 - b) The presence of the second, hidden file system should not be detectable.
- 5. Be evaluated for perceivable video impact using an evaluation study:
 - a) A developed web application should evaluate the requirement “Embedding should occur with no perceivable impact on video quality.” by obtaining data from multiple users.

2.4 Choice of Languages and Tools

With the above requirements for the final product defined, I identified an appropriate set of programming languages and tools.

Stegasis will be designed to operate on the **Linux** operating system since **Windows** has no equivalent of the file system in user space paradigm.

Several of the core (and extension) requirements strongly suggest a lower level language such as C or C++ rather than a higher level sandboxed language such as Java. For example, the parsing and modification of **AVI** files lends itself to a language like C since it will involve large amounts of byte level manipulation. Furthermore, the Microsoft file format reference defines the different data structures used within **AVIs** as C **structs**. **FUSE**, **libjpeg** and libraries provided by **FFmpeg** are all natively C libraries. Although wrappers for other languages (including Java) do exist^{[28] [29]}, they seem to be lacking documentation and few are being actively maintained. The requirement that **Stegasis** should support several steganographic embedding algorithms implores the use of object oriented techniques; defining a **SteganographicAlgorithm** interface of which each embedding algorithm implements. This suggests C++ over C. The final core requirement, performance, also favours C/C++ over Java¹⁴ due to the JVM overheads.

The reasons above and the fact I have prior experience using C++ led to the conclusion that C++ should be the primary language used to develop **Stegasis**.

As discussed in Section 2.1.5, **FFmpeg** will be used for the multiple video formats extension task to allow the decoding and conversion of the many video formats available today. The library **libjpeg** will be used for the manipulation of **JPEG** images.

During the implementation of **Stegasis**, a number of small steganalysis programs will be developed. These will likely be written in a scripting language such as Python or Matlab since both have extensive library support for mathematical operations.

The extension task “**Stegasis** should be evaluated for perceivable video impact using a web application” will require a web application to be developed and hosted for easy access to participants and a database to store the collected user data. **Node.js** together with

¹⁴Numerous studies have shown that C/C++ code performs better than equivalent Java code^{[30] [31] [32]}.

the database `MongoDB` was chosen as the development stack for the site. This decision was mainly due to the speed at which you can develop *CRUD* (create, read, update and delete) web applications and my previous experience with the technologies.

3 || Implementation

3.1 Introduction

The development of **Stegasis** consisted of the five main stages detailed within this chapter. Firstly, a parser for the AVI file format was developed allowing direct access to video pixel data. Next, steganographic embedding algorithms were implemented along with corresponding steganalysis tools to test the security of the proposed techniques. The file system was then developed utilising the AVI decoder and steganographic algorithms to embed and extract data directly into and out of video files. Finally, the extension tasks were individually addressed providing support for multiple video formats, directory structures and plausible deniability.¹⁵ The testing section provides an overview of the testing processes applied throughout development.

The software development process embraced the modern “Launch early, iterate often” methodology^[33]. A simplified version of **Stegasis** was initially produced allowing integration issues to be identified early on, when the code was still very malleable. Once this basic version was working, an iterative approach was then taken to add more functionality and features.

3.2 AVI Decoder

The concept of an AVI decoder is first abstracted to that of a generic **VideoDecoder** interface¹⁶. The core requirements state that the AVI decoder should allow access to individual pixel data. The pixel data within an AVI file is grouped into chunks, one per video frame. This lead to the decision to define the **VideoDecoder** to allow access to the video pixel data at the granularity of a single video frame. It will also be useful for the **VideoDecoder** interface to expose metadata about the video, for example, the total number of video frames in the video, the height and width of the video frames and the total size (in bytes) of each video frame.

This gives the definition for the **VideoDecoder** interface as described in Listing 3.1 (**NextFrameOffset** will be discussed in Section 3.4).

Note that **getFrame** returns a **Frame** wrapper object, rather than a raw **char** pointer to the frame pixel data, adhering to the *Dependency Inversion* principle^[34]. This will be useful when dealing with different video formats that don’t necessarily group all of a frames video data to be accessible by a single **char** pointer.

A **Frame** abstracts the concept of a single frames video data. In the case of uncompressed AVI, this can be thought of as a **char** pointer to the GBR pixel data, along with an associated frame size in bytes. A boolean value is also associated with each chunk, signifying if the chunk data has been modified, that is, it is dirty. The **Frame** interface is defined in Listing 3.2.

¹⁵The evaluation site extension task is discussed within the evaluation chapter.

¹⁶The term “interface” is used as shorthand for an abstract C++ base class. That is, a class with pure virtual member functions and no function implementations.

```

class VideoDecoder {
public:
    virtual Frame *getFrame(int frame) = 0;
    virtual int  getFileSize() = 0;
    virtual int  getNumberOfFrames() = 0;
    virtual int  getFrameSize() = 0;
    virtual int  getFrameHeight() = 0;
    virtual int  getFrameWidth() = 0;

    virtual void getCapacity(char capacity) = 0;
    virtual void setCapacity(char capacity) = 0;

    virtual void getNextFrameOffset(int *frame, int *offset) = 0;
    virtual void setNextFrameOffset(int frame, int offset) = 0;

    virtual void writeBack() = 0;
    virtual ~VideoDecoder() {};
};

```

Listing 3.1: VideoDecoder interface (video/video_decoder.h:15)

```

class Frame {
protected:
    long frameSize;
public:
    virtual long  getFrameSize() = 0;
    virtual char *getFrameData(int n=0, int c=0) = 0;

    virtual bool isDirty() = 0;
    virtual void setDirty() = 0;
};

```

Listing 3.2: Frame interface (video/video_decoder.h:4)

The AVI parsing process can be thought of consisting of two main parts; parsing the video headers and parsing the video chunk data. The pseudocode in Algorithm 3.1 illustrates this with the headers being parsed lines 1-12 and the chunks being parsed lines 15-22.

The actual implementation is slightly more complex than presented above due to the existence of JUNK chunks. The AVI file format specifies that any number of chunks with a **FourCC** code of JUNK and of arbitrary length can be inserted between any AVI list structures. The parser must therefore be able to cope with this.

The **writeBack** function of the AVI decoder will write any modified pixel data back into the original AVI file and **getFrameSize** returns the number of bytes within each frame that should be embedded within, as shown in Listing 3.3.

This expression arises from the fact that uncompressed AVI uses 3 bytes per pixel value. *Capacity* is a user provided percentage ranging in value from 1 - 100. It specifies the percentage of the frame to embed within. **frameSize** must therefore reduce the returned

Algorithm 3.1 AVI parsing process

```

1:  $f \leftarrow \text{open}(\text{file\_path})$ 
2:  $\text{riff\_header} \leftarrow \text{readRiffHeader}(f)$ 
3: if  $\text{riff\_header.fourCC} \neq \text{RIFF}$  then
4:   Error “File is not an AVI file”
5:  $\text{avi\_header} \leftarrow \text{readAviHeader}(f)$ 
6:  $\text{bitmap\_info\_header} \leftarrow \text{readBitmapInfoHeader}(f)$ 
7: if  $\text{bitmap\_info\_header.compression} \neq 0$  then
8:   Error “Stegasis does not natively support compressed AVI files”
9:  $\text{audio\_info\_header} \leftarrow \text{readAudioInfoHeader}(f)$ 
10:  $\text{frame\_chunks} \leftarrow []$ 
11:  $i \leftarrow 0$   $\triangleright$  File pointer is now positioned at the start of the audio video chunks
12: while  $i < \text{avi\_header.total\_frames}$  do
13:    $\text{chunk} \leftarrow \text{readChunk}(f)$ 
14:   if chunk is a video chunk then
15:      $\text{frame\_chunks}[i].\text{chunkSize} = \text{chunk.chunkSize}$ 
16:      $\text{frame\_chunks}[i].\text{frameData} = \text{readChunkData}(f)$ 
17:      $i++$ 
18:   else
19:     Advance  $f$   $\text{chunk.chunkSize}$  bytes

```

```

virtual int frameSize() {
  return (int) floor( this -> aviHeader.width * this -> aviHeader.height * 3 * (
    this -> capacity / 100.0 ));
};

```


Listing 3.3: AVI decoder `frameSize` function (`video/avi_decoder.cc:298`)

frame size value by “capacity percent”.

The effect of the capacity parameter is illustrated within Figure 3.1. The top left image is the original video frame and the top right image is the LSB plane (red channel) of the frame with no data embedded. The bottom two images have data sequentially embedded within them using capacity settings of 50% and 15% respectively.

The image used throughout this chapter for illustration purposes has been specifically chosen to emphasise the effect of the embedding process. Due to post processing applied to the video, the LSB plane closely resembles that of the original frame. This would not be the case for a video recorded on a mobile phone for example.

The AVI decoder as described now provides all necessary functionality to allow the modification of pixel data to achieve the steganographic embedding of information within video frames.



images/cap_ill.png

Figure 3.1: Illustration of the capacity parameter.

3.3 Steganographic Algorithms

As defined within the previous section, a steganographic system consists of a pair of functions providing embedding and extraction functionality. Therefore, a generic `SteganographicAlgorithm` interface will need to declare two functions `embed` and `extract` which will embed and extract data into and out of cover objects. Listing 3.4 shows the interface.

```
class SteganographicAlgorithm {
protected:
    string password;
    VideoDecoder *dec;
public:
    virtual pair<int, int> embed(Frame *f, char *data, int reqByteCount,
        int offset) = 0;
    virtual pair<int, int> extract(Frame *f, char *output, int reqByteCount,
        int offset) = 0;
};
```

Listing 3.4: `StegoAlgorithm` interface (`steg/steganographic_algorithm.h:8`)

The `embed` function attempts to embed `reqByteCount` bytes from `data` into frame `f` starting at an offset `offset` bytes into the frame and returns the number of bytes successfully embedded along with the reached frame offset. Similarly, the `extract` function attempts to extract `reqByteCount` bytes from frame `f` starting at an offset `offset` bytes into the frame and put them into `output` and returns the number of bytes successfully extracted along with the reached frame offset.

3.3.1 LSB Sequential Embedding

Sequential LSB embedding is arguably the simplest steganographic algorithm. It works by replacing the LSBs of the cover object with the bits comprising the message, producing the stego image. Algorithm 3.2 shows pseudocode for the LSB embedding algorithm. The matching extraction algorithm is shown in Algorithm 3.3.

Algorithm 3.2 LSB embedding algorithm

```

1: for  $i \leftarrow 0$  upto dataBytes - 1 do
2:   for  $j \leftarrow 7$  downto 0 do
3:     if The  $j$ th significant bit of data[ $i$ ] == 1 then
4:       Set LSB(frame[offset++]) to 1
5:     else
6:       Set LSB(frame[offset++]) to 0

```

Algorithm 3.3 LSB extraction algorithm

```

1: for  $i \leftarrow 0$  upto dataBytes - 1 do
2:   for  $j \leftarrow 7$  downto 0 do
3:     Set the  $j$ th significant bit of output[ $i$ ] to LSB(frame[offset++])

```

Listing 3.5 shows the actual implementation and Figure 3.2 illustrates the algorithms operation.

```

virtual pair<int, int> embed(Frame *f, char *data, int reqByteCount, int
    offset) {
    this->crypt->encrypt(data, reqByteCount);
    char *frame = f->getFrameData();
    int bytesEmbedded = 0;
    while (bytesEmbedded < reqByteCount && offset < this->dec->getFrameSize()) {
        for (int j = 7; j >= 0; j --) {
            frame[offset] &= ~1;
            frame[offset++] |= ((1 << j) & data[bytesEmbedded]) >> j;
        }
        bytesEmbedded ++;
    }
    this->crypt->decrypt(data, reqByteCount);
    if (offset == this->dec->getFrameSize()) {
        return make_pair(bytesEmbedded, 0);
    } else {
        return make_pair(bytesEmbedded, offset);
    }
};

```

Listing 3.5: LSB embedding implementation (steg/lsb_algorithm.cc:8)

Using the LSB embedding algorithm, 43.2 kB of data is embedded into a 1280×720 video frame (using a capacity setting of 100%). The resulting stego image is shown in Figure 3.3, with the left image the cover object and the right image the stego object.




Figure 3.2: Illustration of the LSB Embedding algorithm, taken from forensicismag.com



Figure 3.3: Effect of the LSB Embedding algorithm.

The visual impact on the video frame is very small and almost certainly not noticeable having been reproduced within this document at a smaller resolution. However, if we take a closer look at a specific portion of the video frame, we can see some small discrepancies between the cover and stego objects - see Figure 3.4. I am unsure how well these images will be reproduced when printed, but the difference is definitely noticeable within the PDF. Without the original cover object for comparison, it would be very hard to identify these details visually and deduce the presence of embedded data. However, if the LSB plane of the frame is visually inspected, as in Figure 3.5, one can immediately detect the presence of the hidden data - the LSB plane of the stego object looks “too random”. This is an example of a *visual steganalysis attack*.

Formalising the looking “too random” idea leads to the *Chi-Squared attack* developed by Westfeld and Pfitzmann^[35]. To implement the Chi-Squared attack, the concept of *Pair of Values* (POVs) is introduced.

One of the effects of an embedding algorithm like sequential LSB embedding is the creation of POVs, pixel values that embed into one another. For example, a pixel value of 100 in the cover image will either stay 100 or change to 101. Similarly, a pixel value of 101 will either stay 101 or change to 100. Thus (100, 101) is a POV.




Figure 3.4: Close up effect of the LSB Embedding algorithm.



Figure 3.5: LSB plane of the cover and stego object.

Definition 3.1. PAIRS OF VALUES

A POV p is a member of the set \mathcal{Pov} , defined as,

$$\mathcal{Pov} \triangleq \{ (2k, 2k+1) \mid 0 \leq k \leq 127 \}$$

Westfeld and Pfitzmann claim that the LSBs in images are not completely random, rather, the frequencies of each of the two pixel values in each POV tend to lie far from the mean of the POV. That is, it is unlikely for the frequency of pixel value $2k$ to be close to equal to the frequency of pixel value $2k+1$. Furthermore, as information is embedded into the cover object, the frequencies of $2k$ and $2k+1$ become (nearly) equal. The Chi-squared attack was designed to detect this and bases the probability of embedding on how close to equal POVs are in the image.

To implement the attack, the following steps are taken in reference to a given image. First, $x_k = \text{frequency}(2k)$ and $y_k = \text{frequency}(2k+1)$ are calculated, followed by the expected frequency $z_k = \frac{x_k + y_k}{2} \forall k$. n is defined to be the number of POVs, $|\mathcal{Pov}|$. For uncompressed AVI using 24 bits per pixel, $n = 128$. The *minimum frequency condition* is now applied. This sets $x_k = y_k = z_k = 0$ and decrements n by one, if the condition $x_k + y_k \leq 4$ holds. The Chi-Squared statistic, with $n-1$ degrees of freedom is then calculated:

$$\chi_{n-1}^2 = \sum_{k=0}^{127} \frac{(x_k - z_k)^2}{z_k}$$

The probability of embedding, p , is then calculated by evaluation of the following integral:

$$p = 1 - \frac{1}{2^{\frac{n-1}{2}} \Gamma(\frac{n-1}{2})} \int_0^{\chi_{n-1}^2} e^{-\frac{u}{2}} u^{\frac{n-1}{2}-1} du$$

See Appendix Section C.1 for an implementation of the Chi-squared attack in Python.

The Chi-squared attack produces very good results for the naïve sequential LSB embedding algorithm. Figure 3.6 show an example of it in use operating on a video frame with a capacity setting of 50%. It obvious that information has been embedded within the first half of the frame.

The Chi-Squared attack motivates the development of a *permuted* LSB embedding algorithm. The attack works so well because regardless of the capacity setting, the data is sequentially embedded from the top of the video frame. It would be better to distribute the data uniformly throughout the entire frame. A permuted LSB embedding algorithm achieves this and resists the Chi-Squared attack when using a small capacity setting.¹⁷

3.3.2 Permuted LSB Embedding

One can produce deceptively simple pseudocode for the permuted LSB embedding algorithm as shown in Algorithm 3.4.

¹⁷A permuted embedding using a capacity of 100% is no different to sequential embedding!

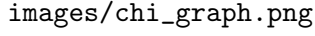


Figure 3.6: Results of the Chi-Squared attack.

Algorithm 3.4 Permuted LSB embedding algorithm

```

1: path ← a pseudorandom permutation of the cover object
2: path.seekToOffset(offset-1)
3: for i ← 0 upto dataBytes - 1 do
4:   for j ← 7 downto 0 do
5:     if The jth significant bit of data[i] == 1 then
6:       Set LSB(frame[path.next()]) to 1
7:     else
8:       Set LSB(frame[path.next()]) to 0

```

Lines 1 and 2 hide a large amount of complexity involved in implementing the algorithm.

Assuming the cover object consists of n bytes, a pseudorandom permutation of the cover object can be thought of as a pseudorandom permutation of the numbers $0 - n$. One way to produce a pseudorandom permutation of a list of numbers is to shuffle an array containing them. This approach has the drawback that you need to hold the numbers in memory (11 MB for a single 720p HD video frame). Instead, I took a different approach using a *Linear Congruential Generator*^[36] (LCG).

Definition 3.2. LINEAR CONGRUENTIAL GENERATOR

A *Linear Congruential Generator* is defined by the recurrence relation:

$$X_{n+1} = (aX_n + c) \bmod m$$

where X is the sequence of pseudorandom values, and

$$\begin{aligned}
 &m, 0 < m - \text{the modulus} \\
 &a, 0 < a < m - \text{the multiplier} \\
 &c, 0 \leq c < m - \text{the increment} \\
 &X_0, 0 \leq X_0 < m - \text{the seed}
 \end{aligned}$$

are integer constants that specify the generator.

The pseudorandom permutation is only needed to uniformly distribute embedding throughout the image and it is therefore not a problem that a LCG may not produce cryptographically secure random numbers. The Hull-Dobell Theorem^[37] states that a LCG will have a full period if and only if the following 3 requirements are satisfied:

1. $\gcd(c, m) = 1$ (c and m are relatively prime),
2. $a - 1$ is divisible by all prime factors of m ,
3. $a - 1$ is a multiple of 4 if m is a multiple of 4.

Therefore, if the above requirements can be satisfied, a LCG can be used as a pseudorandom permutation for the cover object with m equal to n (the size of the cover object). Note that forcing m to be a power of 2 simplifies the above requirements to:

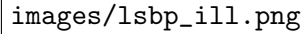
1. $\gcd(c, m) = 1$,
2. $a - 1$ is odd,
3. $a - 1$ is a multiple of 4 if m is a multiple of 4.

As such, m is set equal to the next power of 2 larger or equal to n . Any produced values larger or equal to n are just discarded, therefore producing a full period of size n as required.

The pseudorandom permutation should be dependant on a user provided passphrase. Therefore, the values of a and c are determined using a key derived from the users passphrase. The popular key derivation function PBKDF2^[38] is used along with the Whirlpool^[39] hash function to generate a 128 byte key pool and the first 4 bytes are taken as an unsigned integer and used to derive c and a . Rather than implement these algorithms myself, I used a popular C++ cryptographic library, **Crypto++**.

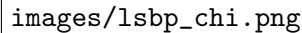
The result of the permuted embedding algorithm is illustrated in Figure 3.7. The left image has data embedded using the sequential embedding algorithm whereas the right image has the same data embedded using the permuted embedding algorithm. Both images were using a capacity setting of 15%.

Figure 3.8 shows the result of applying the Chi-Squared attack to the video frames in Figure 3.7. The left and right graphs show the probability of embedding within the left and right frame respectively. It is clear that the implemented permuted LSB embedding algorithm resists the Chi-Squared attack when using small values for the capacity setting. The above example embedded data within an uncompressed AVI video file which was 10 seconds in duration, had a resolution of 1280×720 and was 741 MB in size. Using the permuted LSB embedding algorithm with a capacity setting of 15% (which resisted the Chi-Squared attack), **Stegasis** informed me that the formatted volume had a total capacity of 14.49 MB, roughly 2% of the file size. This embedding capacity is already a lot better than some of the those provided by programs investigated in Section 2.2, and will be further improved upon during the extension task within Section 3.6.1.



images/lsbp_ill.png

Figure 3.7: Illustration of the permuted LSB algorithm.



images/lsbp_chi.png

Figure 3.8: Chi-Squared attack on the permuted LSB algorithm.

3.3.3 Encrypting the embedded data

To further strengthen the security of **Stegasis**, I decided to offer the option to encrypt data before embedding it within the video. This way, even if the steganography is broken, the information itself will not be compromised. The available cryptographic algorithms are decoupled from the steganographic systems allowing the user to mix and match between them. In addition, the same user provided passphrase used to permute the data throughout the video frame is used within the encryption process¹⁸, addressing another concern of some of the investigated programs in Section 2.2.

AES, **TwoFish** and **Serpent** (all 256 bit) are all available within the **Crypto++** library which made incorporating them in **Stegasis** easy. Counter mode (**CTR**) is used as the block cipher mode so that data is not required to be padded to the block size. A hybrid chained encryption scheme **AES** \rightarrow **TwoFish** \rightarrow **Serpent** is also supported. This method has the advantage that if even 2 of the above algorithms are cryptographically broken,

¹⁸A different section of the 128 bytes key pool is used.

the data will still be secure¹⁹. You do however incur a performance penalty due to the encryption and decryption processes - see Section 4.3.

This concludes the steganographic embedding algorithms developed to satisfy the core requirements. Additional steganographic systems are investigated within the extension tasks along with more appropriate steganalysis methods. Table 3.1 shows a complete list of the implemented steganographic and cryptographic algorithms, including those implemented within the extension tasks (in italics).

Steganographic (Algorithm code)	Cryptographic
LSB Embedding (LSB)	AES
Permuted LSB embedding (LSBP)	Serpent
<i>DCT LSB embedding (DCT)</i>	TwoFish
<i>Permuted DCT LSB embedding (OutGuess 0.1) (DCTP)</i>	AES(Serpent(TwoFish))
<i>F4 (F4)</i>	
<i>F5 (F5)</i>	

Table 3.1: Implemented steganographic and cryptographic algorithms

3.4 The file system

I decided early on that the file system would only support a subset of the features offered by a fully-fledged standard file system. The following features were considered essential for a bare-bones practical file system: Creating and deleting files, reading and writing to files, listing the files in the file system and renaming (moving) files. Support for directories was also considered important and was therefore implemented as an extension task.

File permissions and access / modification times were considered non-essential and therefore not implemented. This means all files will be given the same permissions of 755 (RWXRW-RW-), and all access / modification times default to 0 *Unix Time* (1 January 1970).

Since the file system functionality was being implemented from scratch, the code written turned out complex and intricate due to the many corner cases encountered during integration testing. A large amount of time was dedicated to testing the implemented file system and tracking down particularly nasty concurrency bugs arising in specific circumstances.

The basic idea behind embedding a file system within a video is to break files written to the volume into many *chunks* and embed these into the video frames. Some kind of header will need to be developed to keep track of the locations of these chunks so that each file can be read back correctly.

¹⁹This encryption option is also offered by TrueCrypt.

3.4.1 Developing the header

The header serves a similar purpose to the *File Allocation Table* used with the **FAT** file systems. It will need to be stored in a known location so that it can be extracted when the video file is mounted. I decided to use the first frame of the video to store the header; referred to as the *File Allocation Frame* (**FAF**). The design of the header went through a number of iterations before arriving at the final version presented here, mainly due to underestimating the number of bits needed to store the file frame numbers and offsets. Figure 3.9 shows the overall structure of the **FAF**.

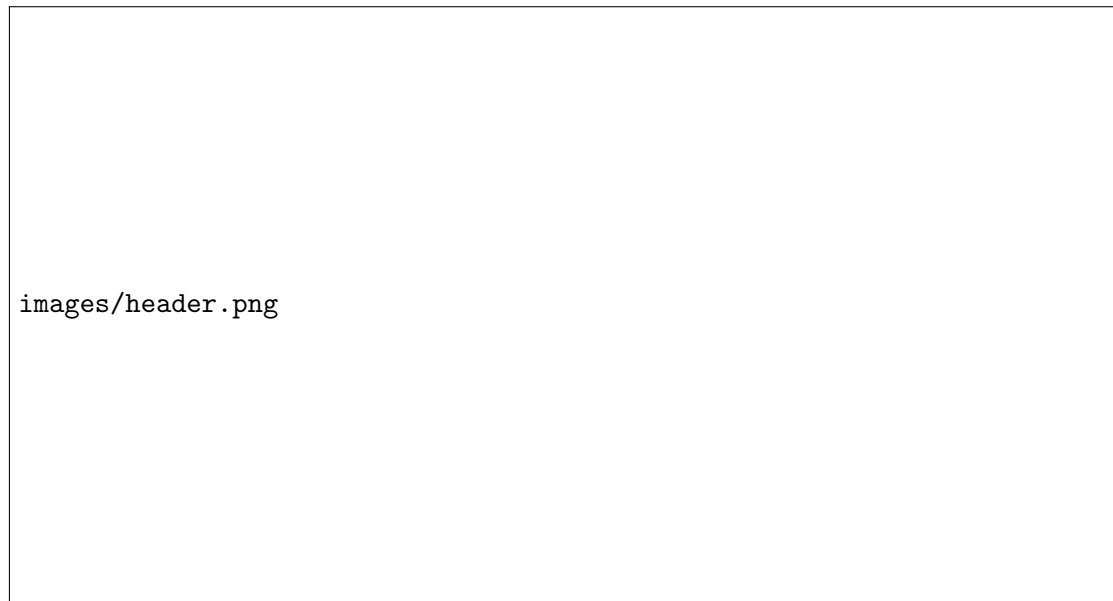


Figure 3.9: File allocation frame structure.

The header section (coloured blue) contains metadata about the embedded file system. “STEG” is the literal 4 character **ASCII** string and is used to check the header has been correctly extracted. If **STEG** is not found as the first 4 bytes of the header, the extraction process is aborted and an error message displayed to the user. The “Header Bytes” field contains the number of remaining bytes to be extracted from the **FAF**. This data contains information for each of the files within the file system, including the file name as a null terminated string, the number of chunks the file is split into and the location and size of each of these chunks. A file chunk is represented as a **FileChunk** struct defined in Listing 3.6.

Within Figure 3.9, a single file chunk is indicated with square brackets, and is coloured green. A file can contain an arbitrary number of chunks each of which can span multiple frames. Note that using an unsigned integer for the bytes field limits individual chunk sizes to roughly 4 GB. However, this will not be a problem since writes to the file system occur in 65 kB chunks.


```

struct FileChunk {
    uint32_t frame;
    uint32_t offset;
    uint32_t bytes;
};

```

Listing 3.6: FileChunk definition (`fs/stegfs.h:19`)

3.4.2 Writing to the file system

A write to the file system occurs in three stages. First, `create` will be called requesting that a new file be created with a specified name. Next, the `write` function will be called a number of times requesting that data be written to the file. Finally `flush` will be called requesting that any data held in memory be flushed to disk.

The first stage is easy to implement and is listed in Listing 3.7.

```

int SteganographicFileSystem::create(const char *path, mode_t mode, struct
    fuse_file_info *fi) {
    this->fileSizes[path] = 0;
    this->fileIndex[path] = std::vector<struct FileChunk>();
    return 0;
};

```

Listing 3.7: The `create` function call (`fs/stegfs.cc:202`).

The write calls are a bit trickier, recall the FUSE write call which requests that `size` bytes from the buffer `buf` should be written to the file `path` starting at offset `offset`. Algorithm 3.5 shows pseudocode for the write function. The full implementation is listed in the Appendix Section C.2.

Algorithm 3.5 Writing to the file system.

```

1: bytesWritten ← 0
2: (nextFrame, nextOffset) ← decoder.getNextFrameOffset()
3: chunk ← FileChunk()
4: while bytesWritten < size do
5:     (bytesEmbedded, newOffset) = embed size-bytesWritten from buf+bytesWritten
        into nextFrame at nextOffset
6:     chunk.bytes += bytesEmbedded
7:     nextOffset = newOffset
8:     if newOffset == 0 then
9:         nextFrame ++
10:    bytesWritten += bytesEmbedded
11: fileIndex[path].append(chunk)

```

Finally, `flush` is called which will write the header and ask the video decoder to write back to disk. This can either occur synchronously or asynchronously depending on a user specified flag.

3.4.3 Reading from the file system

The read function call is similar, recall the declaration which requests that **size** bytes of the file **path** starting at offset **offset** should be written to the buffer **buf**. The read function must identify the chunk **offset** points to and then return the correct amount of data possibly spread across multiple subsequent chunks. Algorithm 3.6 describes the function operation. See the Appendix Section C.2 for the full implementation.

Algorithm 3.6 Reading from the file system.

```
1: bytesWritten  $\leftarrow$  0, chunkNum  $\leftarrow$  0, bytesRead  $\leftarrow$  0
2: for chunk in fileChunks do
3:   if bytesRead + chunk.bytes > offset then
4:     break
5:   else
6:     bytesRead += chunk.bytes
7:     chunkNum ++
8: while bytesWritten < size do
9:   chunk  $\leftarrow$  fileChunks[chunkNum]
10:  bytesLeftInChunk  $\leftarrow$  chunk.bytes - chunkOffset
11:  bytesLeftInChunk  $\leftarrow$  min(bytesLeftInChunk, size-bytesWritten)
12:  extract bytesLeftInChunk from chunk.frame at chunk.offset+chunkOffset into
    buf+bytesWritten
13:  bytesWritten += bytesLeftInChunk
14:  chunkOffset  $\leftarrow$  0
15:  chunkNum ++
```

3.4.4 Listing files in the file system

The function **readdir** is called when the contents of a directory are requested. The implementation is straightforward and just iterates over the **fileSizes** map as shown in Listing 3.8.

```
int SteganographicFileSystem::readdir(const char *path, void *buf,
    fuse_fill_dir_t filler, off_t offset, struct fuse_file_info *fi) {
    filler(buf, '.', NULL, 0);
    filler(buf, '..', NULL, 0);
    for (auto kv : this->fileSizes) {
        filler(buf, kv.first.c_str() + 1, NULL, 0);
    }
    return 0;
};
```

Listing 3.8: FUSE readdir implementation (**fs/stegfs.cc:264**).

The described **FUSE** operations cover the majority of the core file system implementation, successfully implementing all of the decided essential features listed at the start of this subsection.

3.5 Command line application

To complete **Stegasis**, the implemented components must be assembled together and a user interface produced. I decided on a command line interface allowing the user to utilise the functionality developed within this project. Listing 3.9 shows the process of formatting and mounting a video using **Stegasis**.

```
$ stegasis format -alg=lsbp -crypt=aes -pass=123 -cap=25 ~/my_video.avi
Formatting video...
Volume size: 30 MB
Format Successful
$ stegasis mount -alg=lsbp -crypt=aes -pass=123 ~/my_video.avi /mnt/video
Video mounted at /mnt/video
```

Listing 3.9: Using **Stegasis** to format and mount a video.

At this point, the user can copy files into **/mnt/video** and they will be automatically embedded into the video. The user can unmount the video by closing the program (for example by pressing Control-C) and **Stegasis** will gracefully exit, writing back any unflushed changes to disk. See the Appendix Section B for more detailed usage information.

Stegasis works exactly how I envisioned it during the project inception. The example usage I gave within the project proposal is exactly reflected within the finished product²⁰.

3.6 Extension Tasks

All of the extension tasks listed within the project proposal were addressed and with the exception of “Hiding data within audio streams”, all of them were successfully implemented. The further extension task “Plausible deniability” was added during the project and also successfully implemented.

3.6.1 Supporting multiple video formats

As discussed, there are several reasons why it would be beneficial for **Stegasis** to operate on video formats other than uncompressed AVI. To accomplish this, rather than develop many video decoders and steganographic systems for multiple video formats, a generic solution was designed and implemented utilising **FFmpeg**.

Regardless of the video format, a video can be thought of as a sequence of video frames played back at a specific frame rate together with an optional audio track. Using **FFmpeg** to extract the frames of a video and implementing steganographic systems operating on sequences of images, **Stegasis** can function seamlessly across multiple video formats. For

²⁰With the exception of pressing Control-C to exit rather than typing **stegasis unmount**.

example, if a user provides a compressed `h.264` video, `FFmpeg` will extract each frame as a `JPEG` along with the audio track. `Stegasis` will then embed and extract data into and out of the extracted `JPEG` images. When the video is unmounted, the modified images will be re-muxed as a *Motion JPEG* stream together with the audio to produce an `MKV` file - the stego object. The `JPEG` file format was chosen for the intermediate frames due to the large number of steganographic algorithms operating on them. Figure 3.10 shows an overview of the approach.




Figure 3.10: Operation on multiple video formats.

This method will work with virtually any video format presented to `Stegasis` due to the extensive codec library of `FFmpeg`.

Although the initial conversion of the provided video to `JPEG` images can be lossy, the reassembly of video frames and extraction thereafter cannot - this would damage the embedded file system. This constraint leads to the following four `FFmpeg` commands shown in Listing 3.10.

The first two commands convert each frame of the video to `JPEG` images and extract any audio. The reassembly command muxes the modified video frames and audio together into a single `MKV` file. The use of `-codec copy` tells `FFmpeg` to literally copy the `JPEGs` into a Motion JPEG (MJPEG) stream. This ensures the frames are not recompressed - preserving the embedded data. The final command uses `-vcodec copy` which tells `FFmpeg` to losslessly extract each `JPEG` image from a MJPEG stream - preserving the already embedded file system.

Using the above commands, I implemented a video decoder performing the described method. When `Stegasis` is run, a check occurs to see whether the provided file is an uncompressed `AVI`. If it is, the native `AVI` decoder developed is used. If not, this video decoder is used, allowing `Stegasis` to operate seamlessly across all video types without

```
// Convert an arbitrary video to JPEG images (Lossy)
ffmpeg -r <fps> -i <path> -qscale:v 1 -f image2 /tmp/stegasis/image-%d.jpg
// Extract video audio
ffmpeg -i <path> /tmp/stegasis/audio.mp3

// Losslessly convert modified JPEG images to an MJPEG stream
ffmpeg -framerate <fps> -i /tmp/stegasis/image-%d.jpg -i /tmp/stegasis/
    audio.mp3 -codec copy -shortest output.mkv

// Losslessly extract JPEGs from a video already formatted by Stegasis
ffmpeg -r <fps> -i <path> -vcodec copy /tmp/stegasis/image-%d.jpg
```

Listing 3.10: FFmpeg extraction and reassembly commands.

the user needing to manually specify which video decoder to use.

The final step of this extension task is to implement embedding algorithms operating on the extracted JPEG images. Due to the design of **Stegasis**, by implementing the new algorithms satisfying the **SteganographicAlgorithm** interface, all of the file system logic will continue to work as expected.

The implemented algorithms all embed data within the least significant bits of the JPEGs discrete cosine transform coefficients. Several versions of the basic algorithm were implemented along with more advanced algorithms such as F4 and F5. Algorithm 3.7 shows pseudocode for the basic embedding algorithm.

Algorithm 3.7 Basic JPEG embedding algorithm.

```
1: for i ← 0 upto data_bytes - 1 do
2:   for j ← 7 down to 0 do
3:     (row, block, coefficient) ← getCoefficientForOffset(offset++)
4:     component = 2           ▷ Components 2 and 3 are the chroma components.
5:     row ← frame.getRow(row, component)
6:     if The jth significant bit of data[i] == 1 then
7:       Set LSB(row[block][coefficient]) to 1
8:     else
9:       Set LSB(row[block][coefficient]) to 0
```

Within the actual implementation, the component to embed within is chosen using the calculation $(co \% 2) + 1$. This uniformly distributes the embedded data bits between the two chroma components, deliberately not touching the luminance component.

At higher capacity settings, this embedding algorithm does begin to produce visual artifacts within the video frames. Figure 3.11 illustrates this when a capacity setting of 100% is used.

However, at lower capacity settings it is not possible to visually differentiate between




Figure 3.11: Embedding artifacts at 100% capacity.

the cover and stego objects²¹. Roughly, any capacity setting above 50% begins to introduce visually detectable artifacts. Figure 3.12 shows the same frame as above, but using a capacity setting of 20%. (Left image is the original frame.)




Figure 3.12: Frame comparison at 20% capacity.

The basic JPEG embedding algorithm and the permuted variant (**OutGuess 0.1**) are both vulnerable to the *histogram attack* discussed in Section 4.2. This motivates the implementation of more advanced algorithms such as **F4** and **F5**^[40] which resist this type of steganalysis attack. Instead of replacing the LSB of DCT coefficients, **F5** decrements the coefficients absolute value via *matrix encoding*. This results in no pairs of values being produced and thus **F5** cannot be detected by the Chi-Squared attack. Matrix encoding within **F5** uses *Hamming codes* to embed a k -bit message into a $2^k - 1$ bit code word, changing it at most by one bit. This allows **F5** to achieve high embedding efficiencies. See the Appendix Section C.3 for more details and an implementation of **F5**.

The above described additions to **Stegasis** allow it to seamlessly operate across a large range of video formats, greatly increasing practicality and successfully completing the extension task.

²¹This claim is verified within the evaluation Section 4.2.

3.6.2 File system directory structures

The core implementation of the file system does not allow the creation and manipulation of directories, forcing all files written to the volume to reside in the root. Although **Stegasis** is still usable with this limitation, it would be nice to allow users to organise their embedded files using directories as you would expect from a normal file system. To achieve this, the **mkdir** FUSE operation will need to be implemented along with a few changes made to the current file system implementation.

A third data structure, **dirs** is first added along side **fileSizes** and **fileIndex** containing each of the directories within the file system. The **mkdir** operation is then trivial to implement as shown in Listing 3.11.

```
int SteganographicFileSystem::mkdir(const char *path, mode_t mode) {
    this->dirs.insert(path);
    return 0;
}
```

Listing 3.11: FUSE mkdir implementation (**fs/stegfs.cc:158**).

The **readdir** implementation will need to be modified to correctly list the files with respect to the current working directory. It is no longer correct to just iterate over all files in the file system and return their names since you only want to return a file if the function call is requesting the directory that file directly resides within. This can be accomplished by testing if the requested path is a prefix of the file name and by checking the number of slashes in the file name is equal to the number of slashes in the path. Algorithm 3.8 describes the **readdir** implementation.

Algorithm 3.8 Algorithm for **readdir**.

```
1: add('.', add('.'))
2: pathSlashes ← number of slashes in path
3: for file in (fileSizes and dirs) do
4:     fileSlashes ← number of slashes in file path
5:     if path == / then
6:         if file path contains one slash then
7:             add(file)
8:     else if path is a prefix of file and pathSlashes == fileSlashes - 1 then
9:         add(file)
```

Directories will also need to be written to the header. I chose to represent a directory as a file in the header of the video which has a value of -1 in the **number of triples** field. This meant the structure of the **FAF** did not need to be modified. The **readHeader** and **writeHeader** functions were modified appropriately.

These changes successfully implement the extension task and greatly improve the usability of the file system. With the exception of permissions and access / modification times, **Stegasis** now provides all the functionality one would expect from a typical file system.

3.6.3 Plausible deniability

Similar to TrueCrypt's hidden volume feature, I planned to implement plausible deniability by embedding two separate file systems within one video, each using a different user provided passphrase. The *outer* volume will reside at the beginning of the video (where it usually would) and the user should populate this volume with files they are willing to reveal if necessary. The *inner* volume will reside half way through the video file and the user should place sensitive files here. When forced to give up the encryption keys, the user can reveal the passphrase for the outer volume in confidence the inner volume will not be compromised.

This method is vulnerable to steganalysis attacks that would be able to detect the presence of the hidden volume half way through the video. To combat this, random data is embedded throughout the entire video during the format process using both passphrases. Since the volumes are encrypted, it will not be possible to tell if the identified embedded data is a hidden volume or just the random data written during the formatting process^[42].

A second passphrase command line flag **pass2** was added, if this is specified during the format process two headers are written at the start and in the middle of the video. The mount process now attempts to extract and decrypt the first header using the provided passphrase. If this fails, it then attempts to decode the second header (embedded within the middle frame). Depending on which header extracts successfully, the outer or hidden file system is presented to the user. Listing 3.12 shows an example use case of the plausible deniability functionality.

```
$ stegasis format -alg=lsbp -crypt=serpent -pass=outer -pass2=hidden
  -cap=20 video.avi
[...]
Format successful!
$ stegasis mount -alg=lsbp -crypt=serpent -pass=outer video.avi /tmp/steg
[...]
Mounting...

[Second Terminal]
$ echo 'outer' > /tmp/steg/outer.txt

[First Terminal]
<Control-C>
[...]
$ stegasis mount -alg=lsbp -crypt=serpent -pass=hidden video.avi /tmp/steg
[...]
Mounting...

[Second Terminal]
$ echo 'inner' > /tmp/steg/sensitivefile.txt
```

Listing 3.12: Stegasis plausible deniability functionality.

Since the header and data for the hidden file system is stored within the middle video frame onwards, it is possible to damage it by writing too much data to the outer volume.

The user has effectively sacrificed half of the total embedding capacity in return for plausible deniability.

3.6.4 Hiding data within audio

I stated within the project proposal that “a substantial part of an AVI file may be the audio data”. I was correct to use the word “may” within this statement since it turns out that a very small percentage of an uncompressed AVI file is the audio data. Similarly low percentages apply to other video formats due to modern audio compression algorithms such as MP3. The question of how this extra embedding capacity could actually be utilised is also quite prevalent since **Stegasis** and the file system logic all rely on the idea of the video being broken up into frames. This concept does not translate easily over to audio. I therefore decided not to pursue this extension task.

3.6.5 Evaluation of embedding impact on video quality

To evaluate the requirement “Embedding should occur with no perceivable impact on video quality.”, a web application was developed and hosted. Figure 3.13 shows a screenshot of the finished website hosted at www.stegasis.co.uk for the majority of the duration of the project.




Figure 3.13: Screen shot of the evaluation study site.

The user is presented with two images and asked the question “Which image do you think contains hidden data?”. The user may inspect the images for as long as they wish, and by clicking on an image, they select their choice.

In total there are 14,912 pairs of images, one of which is randomly selected for each page load. One of the images is the original video frame while the other has 2.5 kB of data embedded within it using the permuted JPEG embedding algorithm and a capacity setting of 20% (the default value). The position of the “correct” image is randomised and the file names of the images are random strings; attempting to thwart any attempts at foul play. When a user first visits the site, a cookie containing a unique ID is issued allowing data points from the same user to be aggregated together.

The results of this evaluation study obtained by the implemented web application are discussed in Section 4.2 and conclude that it was not possible for the users to visually differentiate between the cover and stego objects at the default capacity setting of 20%.

3.7 Testing

Stegasis was tested using a combination of unit, integration and visual testing. Due to the decoupled nature of the steganographic embedding algorithms, it was easy to produce unit tests for the embedding and extraction functionality. A test suite was written testing all stenographic, cryptographic algorithms and associated code. Listing 3.13 gives an example unit test.

```
void AlgorithmShouldCorrectlyExtractEmbeddedData() {
    char *randomData = (char *)malloc(sizeof(char) * 10);
    for (int i = 0; i < 10; i++) { randomData[i] = rand(); }
    pair<int, int> output = this->alg->embed(this->dec->getFrame(0),
        randomData, 10, 0);
    ASSERT(output.first == 10, "StegAlgorithm did not embed 10 bytes as
        requested\n")

    char *outputData = (char *)malloc(sizeof(char) * 10);
    this->alg->extract(this->dec->getFrame(0), outputData, 10, 0);
    for (int i = 0; i < 10; i++) {
        ASSERT(randomData[i] == outputData[i], "StegAlgorithm extracted data
            did not match requested embedded data\n")
    }
};
```

Listing 3.13: Embedding algorithm unit test (`steg/steg_algorithm_test.cc:83`).

The above unit test verifies the steganographic correctness of the stego systems, checking that data embedded is correctly extracted after being embedded. See the Appendix Section D.1 for the complete test suite.

I deemed it more appropriate to test the file system functionality via an integration test suite due to its highly coupled nature. The test suite revolved around a number of compressed archives containing test files and directory structures. **Stegasis** is first instructed to mount a test video into `/tmp/test`. The test archives are then decompressed into the file system which is then traversed to check the contents is as expected.

Visual testing was also employed. For example, large media content such as high definition video was copied into the mounted volume and checked to see if it played back correctly. This is a good test case since the file (due to its size) will be spread across a large proportion of the video. Indeed, several bugs were identified using this testing approach.

The AVI parser was tested using a black box testing approach on a number of different uncompressed AVI videos. The parser prints out debug information about the video file which can be visually inspected. Checking that the AVI still plays back after steganographic modification is also a good indication that the parser is operating successfully.

See the Appendix Section D for some more testing code samples.

4 || Evaluation

4.1 Satisfaction of Requirements

Stegasis satisfies all of the core requirements and addresses all of the proposed extension tasks. I therefore consider the project a success. Each requirement will now be addressed in turn.

4.1.1 Embedding data within video files

Stegasis does indeed allow data to be embedded within video files. Furthermore, it provides a total of 6 different steganographic embedding algorithms all of which satisfy correctness²² and, an evaluation study discussed in Section 4.2 below concludes that **Stegasis** can operate with no perceivable impact on video quality. A number of steganalysis tools were developed to test the security of the proposed steganographic systems, prompting the design and implementation of more secure embedding algorithms. Encryption functionality is also provided making use of the user provided passphrase. Finally, a capacity setting was also implemented allowing users to choose a trade off between steganographic security and embedding capacity.

4.1.2 Providing a file system interface

Stegasis allows users to specify a mount point at which a logical volume will be presented. Data written to and read from this volume will be embedded into and extracted from the video on the fly. Standard Unix tools including `cp`, `mv` and `rm` work as expected along with directory structures. The file system is correctly persevered between unmounts and remounts of the same video.

4.1.3 Supporting uncompressed AVI

Raw uncompressed AVI files are supported via a developed native parser which allows access to and modification of individual pixel data along with extraction of meta-data about the video.

4.1.4 File system performance

The performance of the file system allows full HD video content to be played back from directly within the volume (12 Mb/s) for all embedding algorithms except F4 and F5. The file system can provide read and write speeds in excess of 20 MB/s (up to 96 MB/s) but this is dependent on the embedding algorithm and encryption scheme chosen. File system performance is discussed in more detail in Section 4.3.

²²Verified by unit test cases as explained in Section 3.7.

4.1.5 Supporting multiple video formats

The novel method described in Section 3.6.1 using FFmpeg allows **Stegasis** to seamlessly operate across virtually every video format available today. The basic implemented embedding algorithms operating on JPEG images (DCT, DCTP) offer large embedding capacities allowing multiple large files to be hidden inside of a single video, while the more advanced implemented algorithms (F4, F5) offer high steganographic security.

4.1.6 File system directory operations

Directory functionality was successfully implemented allowing the creation and manipulation of directory structures, enabling users to organise their embedded files within folders. The `mv` and `rm` commands work as expected when operating on directories as well as normal files.

4.1.7 Embedding within audio

This extension task was investigated but I decided not to pursue it due to the relatively little embedding capacity gain compared to the work required to implement it.

4.1.8 Plausible deniability

Plausible deniability functionality was implemented allowing a second hidden file system to be optionally embedded within a video. Depending on the passphrase provided by the user during mounting, either the outer or inner volume is presented. This satisfies the first requirement point. The second point, that the hidden volume should not be detectable, has also been satisfied since it is not possible to differentiate between the random data written during the format process and the second encrypted hidden volume^[42].

4.1.9 Evaluation of the visual impact of embedding

A web application was implemented and hosted for the majority of the duration of the project. A total of 2040 data points were collected from 21 unique users. The results are discussed below in Section 4.2, and the outcome of the study confirmed that it was not possible for the users to visually perceive embedding artifacts produced by **Stegasis** when using the default capacity setting (20%). This satisfies the final extension task.

Stegasis therefore has addressed all of the proposed extension tasks, further cementing the success of the project.

4.2 Security

Although security of the steganographic algorithms was not a major focus of this project, steganalysis techniques were implemented and these led to more secure algorithms being developed. The user study was also focused on security - attempting to decide if it is possible to visually differentiate between cover and stego objects produced by **Stegasis**

when using its default capacity setting. The raw data obtained from the user study is included in the Appendix Section E and is now statistically analysed.

Due to the layout of the website, if a user cannot tell the difference between the two presented images, they will select one arbitrarily. From the phrasing of the question posed to the user, a response is considered correct if the image selected contained hidden data and incorrect if it did not. If a user were to randomly guess each time, the resulting data stream would be a random stream of corrects and incorrects. Letting correct be represented by 1 and incorrect by 0, this stream becomes a random bit string. Therefore, if it can be shown that the collected data is a random bit string, it can then be concluded that the users must have been randomly guessing and therefore could not differentiate between the two presented images.

To determine if the obtained bit string is in fact random, firstly a hypothesis test is constructed to test if the bits are i.i.d. Bernoulli(1/2) which would be the case for a truly random bit string.

Let the bit string $\mathbf{X} = X_1 \dots X_n$, with each $X_i \stackrel{iid}{\sim} \text{Bernoulli}(p)$ and define $T = \sum_{i=1}^n X_i$. T is a sufficient statistic for p and clearly $T \sim \text{Binomial}(n, p)$.

Test $H_0 : p = \frac{1}{2}$ against $H_1 : p \neq \frac{1}{2}$ at significance level α . Under H_0 , $T \sim \text{Binomial}(n, \frac{1}{2})$, by the de Moivre-Laplace theorem $T \sim N(\frac{n}{2}, \frac{n}{4})$ for large n . The $(1 - \alpha)$ confidence interval for T therefore is

$$\left[\frac{n}{2} - z_{\alpha/2} \sqrt{\frac{n}{4}}, \frac{n}{2} + z_{\alpha/2} \sqrt{\frac{n}{4}} \right].$$

For the bit string observed from the user study, $n = 2040$ and $T = 1042$. At $\alpha = 5\%$, $z_{\alpha/2} = 1.9600$ and the confidence interval is $[975, 1065]$ which contains T . Therefore there is insufficient evidence to reject the null hypothesis and therefore insufficient evidence to conclude the bit string is not random.

Secondly, the National Institute of Standards and Technology Statistical Test Suite^[43] is used to evaluate the bit string. The results are summarised in Listing 4.1.

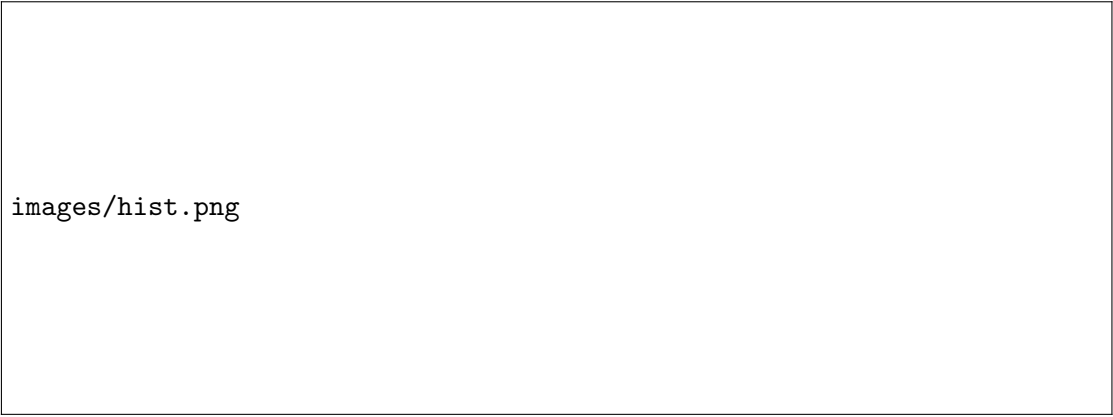
C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	P-VALUE	PROPORTION	TEST
6	1	5	18	0	3	0	5	0	2	0.000000	* 40/40	Frequency
0	0	0	0	0	0	0	0	0	40	0.000000	* 40/40	BlockFrequency
3	4	1	2	3	13	4	5	3	2	0.002465	40/40	CumulativeSums
3	3	0	8	4	11	0	6	1	4	0.000954	40/40	CumulativeSums
5	11	4	4	4	2	5	1	2	2	0.035174	38/40	Runs
40	0	0	0	0	0	0	0	0	0	0.000000	* 0/40 *	LongestRun
2	7	0	17	0	0	0	14	0	0	0.000000	* 40/40	FFT
0	0	0	0	0	0	0	0	0	40	0.000000	* 40/40	ApproxEntropy
1	0	0	0	39	0	0	0	0	0	0.000000	* 39/40	Serial
0	0	0	0	40	0	0	0	0	0	0.000000	* 40/40	Serial
15	0	0	3	0	0	0	0	7	15	0.000000	* 37/40	LinearComplexity

The minimum pass rate for each statistical test is approximately = 37 for a sample size = 40 binary sequences.

Listing 4.1: Output of the NIST Statistical Test Suite.

With the exception of the Longest Runs test (which I believe has failed due to the small sample size), the bit string has exceeded the minimum pass rate for each statistical test. These results strongly indicate that the bit string is random and therefore that the users could not differentiate between the cover and stego objects produced by **Stegasis**.

Although we have concluded that **Stegasis** can operate with no perceivable visual impact on video quality, this does not imply that it will resist further steganalysis techniques. For example, the *Histogram attack* which operates on JPEG images breaks the basic embedding algorithms (DCT and DCTP) implemented in the multiple video formats extension task (even using low capacity settings). The histogram attack enumerates the frequencies of the DCT coefficients of a given image. A characteristic of unmodified JPEG images is that its histogram is symmetrical about 0. This characteristic is lost when embedding occurs. Figure 4.1 shows the effectiveness of the attack, comparing the histogram of an original video frame to one with embedded data. It clearly manages to differentiate between the cover and stego objects.



images/hist.png

Figure 4.1: The histogram attack on the DCTP algorithm.

F4 and F5 however resist this attack and provide a much higher level of steganographic security. Jessica Fridrich et al. showed they were able to break the F5 algorithm by estimating the cover image histogram by slightly cropping the stego image.^[41] However, **Stegasis** is able to spread the embedded data across thousands of JPEG images using a small capacity setting, so I am unsure how effective this attack would be.

The framework **Stegasis** provides makes it easy to implement more secure embedding algorithms which would be a possible future project direction.

4.3 Performance

The performance of **Stegasis**, and in particular the file system, is very important to consider. As with **TrueCrypt**, using **Stegasis** should constitute a similar experience to copying files onto removable storage.

4.3.1 Performance of the file system

Due to the large embedding capacities offered, large files are expected to be copied into the file system and embedded into the video. The write performance of the file system was therefore given a lot of attention during the implementation. The biggest improvement in file system performance was obtained by enabling the FUSE `big_writes` mount option. By default, FUSE limits write calls to blocks of 4096 bytes, `big_writes` allows this limit to be increased to 65 kB. Figure 4.2 shows the effect of varying the block size on write performance for the LSB embedding algorithm.

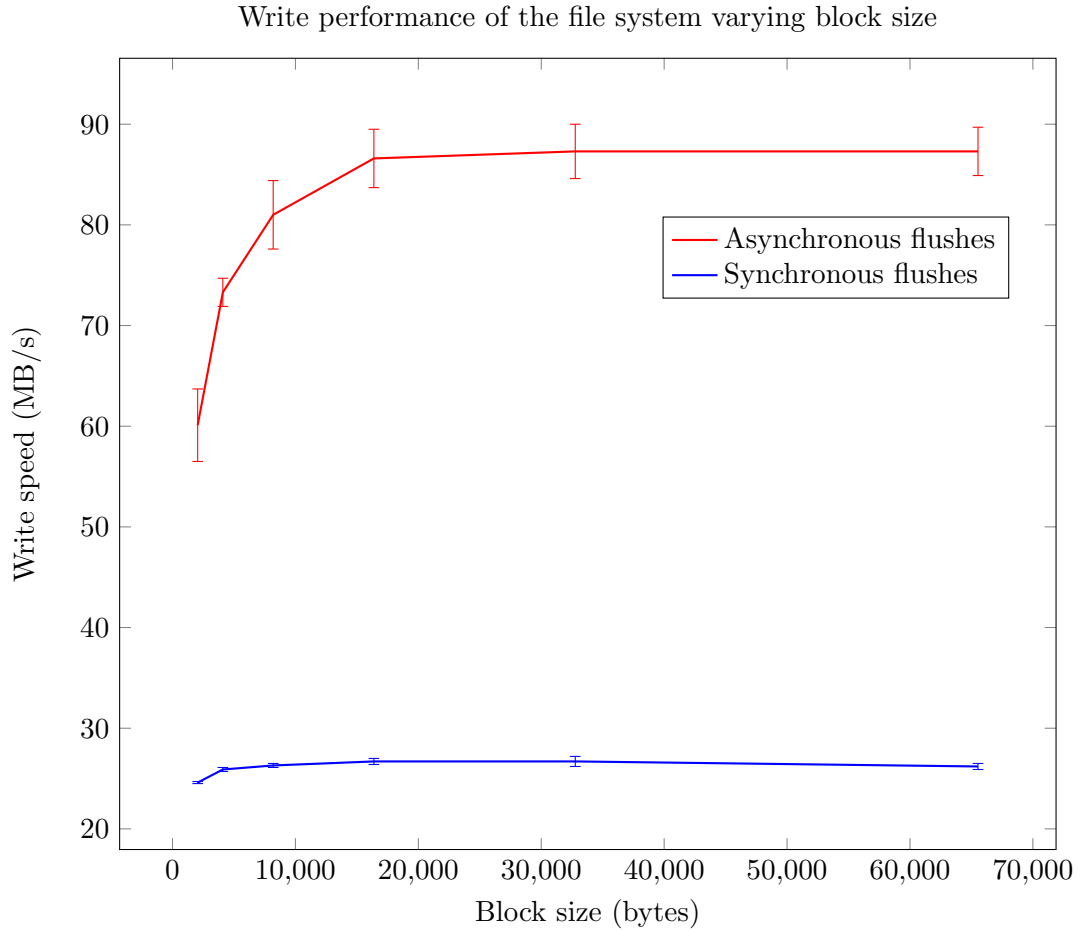


Figure 4.2: Write performance of the file system.

To further evaluate the performance of the file system, I timed the copying of random data into the file system for each combination of steganographic and cryptographic algorithm (30 in total) using the `time` tool. The results of this are shown in Table 4.1. As expected, the embedding algorithms which permute and encrypt data give worse performance, but I think this is a fair trade off.

Stego System	Crypto System	File Size	Copy time (s)	Speed (MB/s)
LSB	<i>None</i>	50 MB	0.521 ± 0.03	96.08 ± 5.10
LSB	AES	50 MB	0.603 ± 0.04	83.02 ± 4.99
LSB	Twofish	50 MB	1.096 ± 0.03	45.63 ± 1.38
LSB	Serpent	50 MB	1.600 ± 0.02	31.24 ± 0.29
LSBP	<i>None</i>	50 MB	7.94 ± 0.15	6.30 ± 0.12
LSBP	AES	50 MB	8.18 ± 0.28	6.11 ± 0.21
LSBP	Twofish	50 MB	8.57 ± 0.38	5.84 ± 0.26
LSBP	Serpent	50 MB	9.24 ± 0.33	5.41 ± 0.20
DCTL	<i>None</i>	25 MB	14.63 ± 0.05	1.71 ± 0.01
DCTL	AES	25 MB	14.69 ± 0.14	1.70 ± 0.02
DCTL	Twofish	25 MB	15.15 ± 0.04	1.65 ± 0.00
DCTL	Serpent	25 MB	15.74 ± 0.01	1.59 ± 0.01
DCTP	<i>None</i>	25 MB	17.29 ± 0.13	1.45 ± 0.01
DCTP	AES	25 MB	17.35 ± 0.13	1.44 ± 0.01
DCTP	Twofish	25 MB	17.75 ± 0.07	1.41 ± 0.01
DCTP	Serpent	25 MB	18.36 ± 0.05	1.36 ± 0.00
F4	<i>None</i>	1 MB	15.40 ± 0.03	0.065 ± 0.00
F4	AES	1 MB	15.54 ± 0.07	0.064 ± 0.00
F4	Twofish	1 MB	17.09 ± 0.09	0.059 ± 0.00
F4	Serpent	1 MB	18.82 ± 0.05	0.053 ± 0.00
F5	<i>None</i>	150 kB	12.07 ± 0.05	0.067 ± 0.00
F5	AES	150 kB	12.35 ± 0.26	0.012 ± 0.00
F5	Twofish	150 kB	12.93 ± 0.04	0.012 ± 0.00
F5	Serpent	150 kB	13.97 ± 0.06	0.011 ± 0.00

Table 4.1: Steganographic and cryptographic algorithm combination performance.

The LSB algorithm gives write speeds on par with SATA hard drives^[44] far exceeding those of USB 2.0 (and even 3.0 devices^[45]) which was the original requirement. The LSBP algorithm is also almost on par with USB 2.0 devices, but the algorithms involving decompressing JPEG images perform a lot worse. This is likely another area for future project work.

5 || Conclusions

Informally, I wanted this project to result in the “**TrueCrypt** of video steganography”, focusing on a practical application allowing multiple files to be easily steganographically embedded within video files utilising a file system interface. This involved researching multiple technical fields (mainly steganography, file systems and various video formats) and developing an application to combine them all into a polished tool.

A number of steganographic embedding algorithms were researched and implemented borrowing ideas and algorithms from cryptography to further increase security. A file system in user space residing within a video file was designed and developed offering the majority of the functionality one would expect from a standard file system. A native uncompressed AVI parser was implemented along with a novel method of supporting other video formats utilising FFmpeg. Finally a Linux command-line application, **Stegasis**, was produced combining the above into a single application.

Stegasis satisfies all of the core project requirements and addresses all of the proposed extension tasks. In reference to the evaluation section of this document, I very much consider this project a success and hope to release **Stegasis** in the near future.

5.1 Lessons Learnt

The “Launch early, iterate often” approach taken to development did come with some drawbacks. Most notably, design decisions made early on during development were sometimes not given as much thought as they possibly should have. This resulted in some issues arising much later during development which required a lot of work to correct. Had due thought been given to these decision at the start of the development process, a large amount of time spent refactoring code could have been saved.

5.2 Future Project Directions

Further work would likely involve implementing more embedding algorithms. Another avenue to explore would be implementing more native video decoders. For example, I am currently working on a native h.264 decoder which embeds data within motion vectors. This has the advantage over the FFmpeg method that the video file would be modified in place as occurs with AVI files (rather than creating a new MKV file). Finally, it would be great for **Stegasis** to be cross platform - developing a Linux only application severely limits the target audience. However, since the file system was implemented using FUSE (which is not compatible with Windows) this may prove tricky.

Overall I believe I have developed a useful application which will hopefully make video steganography more accessible to users. Steganographically embedding file systems within video files has as far as I can tell not been researched as of yet and **Stegasis** is the first available program to support this functionality.

References

- [1] *Steganography in Digital Media*, Jessica Fridrich 2010, Cambridge University Press (Section 1)
- [2] *Regulation of Investigatory Powers Act 2000*
<http://www.legislation.gov.uk/ukpga/2000/23/section/53> (Section 1)
- [3] *A Survey of Steganographic and Steganalytic Tools for the Digital Forensic Investigator*, Pedram Hayati et al.
http://www.pedramhayati.com/images/docs/survey_of_steganography_and_steganalytic_tools.pdf (Section 1)
- [4] *TrueCrypt is a source-available freeware utility used for on-the-fly encryption*.
<http://www.truecrypt.org> (Section 1)
- [5] *The NSA Thinks You Are an Extremist If You Care About Privacy*, Fahmida Rashid 2014, PCMag
<http://securitywatch.pcmag.com/privacy/325273-the-nsa> (Section 1.1)
- [6] *The prisoners' problem and the subliminal channel*, Gustavus Simmons 1984, Advances in Cryptology, pp. 51 - 67 (Section 2.1)
- [7] *La Cryptographie Militaire*, Auguste Kerckhoffs 1883, Journal des sciences militaires, vol. IX (Section 2.1)
- [8] *Steganalysis: The Investigation of Hidden Information*, Neil Johnson and Sushil Jajodia 1998, Information Technology Conference (Section 2.1.2)
- [9] *Using Video Data*, ATLAS.ti 6 2011
http://atlasti.com/wp-content/uploads/2014/05/video_formats.pdf (Section 2.1.3)
- [10] *AVI RIFF File Reference*, Microsoft
<https://msdn.microsoft.com/en-us/library/windows/desktop/dd318189%28v=vs.85%29.aspx> (Section 2.1.3)
- [11] *Steganography and Steganalysis of JPEG Images*, Mahendra Kumar
<https://www.cise.ufl.edu/~makumar/proposalppt.pdf> (Section 2.1.4)
- [12] *JSteg: Steganography and Steganalysis*, Murali P 2009
<http://csis.bits-pilani.ac.in/faculty/murali/netsec-09/seminar/refs/anuroopsrep.pdf> (Section 2.1.4)
- [13] *Hide and Seek: An Introduction to Steganography*, Niels Provos and Peter Honeyman 2003, p. 36
<http://niels.xtdnet.nl/papers/practical.pdf> (Section 2.1.4)

- [14] *Libjpeg, a widely used C library for reading and writing JPEG image files*
<http://libjpeg.sourceforge.net> (Section 2.1.4)
- [15] *Human eye sensitivity and photometric quantities*
<https://www.ecse.rpi.edu/~schubert/Light-Emitting-Diodes-dot-org/Sample-Chapter.pdf> (Section 2.1.4)
- [16] *FFmpeg: A complete, cross-platform solution to record, convert and stream audio and video*
<https://www.ffmpeg.org/> (Section 2.1.5)
- [17] *FFmpeg Codecs Documentation*
<https://www.ffmpeg.org/ffmpeg-codecs.html> (Section 2.1.5)
- [18] *FUSE: File system in userspace*
<http://fuse.sourceforge.net> (Section 2.1.6)
- [19] *Linus vs FUSE: Kernel file system vs FUSE*, Sage 2011
<https://ceph.com/dev-notes/linus-vs-fuse> (Section 2.1.6)
- [20] *Stegostick: A steganographic Tool that lets you hide any file into any file*, 2008
<http://sourceforge.net/projects/stegostick> (Section 2.2.1)
- [21] *Stegomagic: An encrypting software application designed specifically for helping you hide files or messages in media items*, 2014
<http://www.softpedia.com/get/Security/Encrypting/StegoMagic.shtml> (Section 2.2.2)
- [22] *Real steganography with TrueCrypt*, Martin Fiedler 2011
<http://keyj.emphy.de/real-steganography-with-truecrypt> (Section 2.2.3)
- [23] *Stegovideo: A unique tool for hiding information in video*, Oleg Petrov 2011
http://www.compression.ru/video/stego_video/index_en.html (Section 2.2.4)
- [24] *OpenPuff: A steganography tool supporting many carrier formats*, EmbeddedSW
http://embeddedsw.net/OpenPuff_Steganography_Home.html (Section 2.2.5)
- [25] *Steganosaurus is a dissertation project exploring the application of video steganographic and video steganalysis techniques.*, James Ridgway 2014
<http://www.steganosaur.us> (Section 2.2.6)
- [26] *YouTube Advanced encoding settings.*, Google 2014
<https://support.google.com/youtube/answer/1722171?hl=en-GB> (Section 4a)
- [27] *Example USB 2.0 flash drive*, Ebuyer 2015
<https://www.ebuyer.com/543103-toshiba-transmemory-8gb> (Section 4b)
- [28] *JavaFUSE provides Java bindings for FUSE*, Aditya Rajgarhia 2010
<https://code.google.com/p/javafuse/> (Section 2.4)

-
- [29] *TurboJPEG: A Java interface for libjpeg-turbo*
<http://www.libjpeg-turbo.org> (Section 2.4)
- [30] *Numeric performance in C, C# and Java*, Peter Sestoft 2010
<https://www.itu.dk/people/sestoft/papers/numericperformance.pdf> (Section 14)
- [31] *Time Comparing between Java and C++ Software*, Asad Mahmoud Alnaser et al. 2012
<http://www.scirp.org/journal/PaperDownload.aspx?paperID=21960> (Section 14)
- [32] *A Java vs. C++ performance evaluation: a 3D modeling benchmark*, L. Gherardi et al.
<http://www.best-of-robotics.org/pages/publications/gherardi12java.pdf> (Section 14)
- [33] *Release Early, Release Often*, Eric S. Raymond
<http://www.catb.org/~esr/writings/cathedral-bazaar/cathedral-bazaar/ar01s04.html> (Section 3.1)
- [34] *The Dependency Inversion Principle*
<http://www.objectmentor.com/resources/articles/dip.pdf> (Section 3.2)
- [35] *Attacks on Steganographic Systems: Breaking the Steganographic Utilities EzStego, Jsteg, Steganos, and S-Tools and Some Lessons Learned*, Andreas Westfeld and Andreas Pfitzmann 2000 (Section 3.3.1)
- [36] *The Linear Congruential Method*, Donald Knuth, *The Art of Computer Programming*, Volume 2, p. 9 (Section 3.3.2)
- [37] *Random number generators*, T E. Hull A. R. Dobell 1962, *SIAM Review*, Volume 4, Issue 3 (Section 3.3.2)
- [38] *PKCS #5: Password-Based Cryptography Standard*, RSA Laboratories, 2000 <https://tools.ietf.org/html/rfc2898> (Section 3.3.2)
- [39] *The Whirlpool Secure Hash Function*, Williams Stallings 2006
https://www.seas.gwu.edu/~poorvi/Courses/CS381_2007/Whirlpool.pdf (Section 3.3.2)
- [40] *F5: steganographic algorithm: High capacity despite better steganalysis*, Andreas Westfeld 2001, 4th International Workshop on Information Hiding
<https://f5-steganography.googlecode.com/svn/tags/release-0.3/F5%20Steganography.pdf> (Section 2.1.4), (Section 3.6.1), (Section C.3)
- [41] *Steganalysis of JPEG Images: Breaking the F5 Algorithm*, Jessica Fridrich et al., 2002, 5th International Workshop
<http://ws2.binghamton.edu/fridrich/Research/f5.pdf> (Section 4.2)

- [42] *Random versus Encrypted Data*, A. G. Basile 2008
<http://opensource.dyc.edu/sites/default/files/random-vs-encrypted.pdf>
(Section 3.6.3), (Section 4.1.8)
- [43] *National Institute of Standards and Technology Statistical Test Suite*, Juan Soto 1999
http://csrc.nist.gov/groups/ST/toolkit/rng/documentation_software.html (Section 4.2)
- [44] *HDD Write Throughput Average Benchmark*, Tom's Hardware 2013
<http://www.tomshardware.co.uk/charts/hdd-charts-2013/-04-Write-Throughput-Average-h2benchw-3.16,2904.html> (Section 4.3.1)
- [45] *Lexar 8GB USB 3.0 JumpDrive*, Ebuyer 2015
<https://www.ebuyer.com/399229-lexar-8gb-jumpdrive> (Section 4.3.1)

A || Details of the AVI file format

A.1 Detailed AVI form

Listing A.1 shows an expanded form of the AVI structure.

```
RIFF (
  'AVI '
  LIST (
    'hdlr'
    'avih'(<Main AVI Header>)
    LIST (
      'strl'
      'strh'(<Stream header>)
      'strf'(<Stream format>)
      [ 'strd'(<Additional header data>) ]
      [ 'strn'(<Stream name>) ]
      ...
    )
    ...
  )
  LIST (
    'movi'
    {SubChunk |
      LIST (
        'rec '
        SubChunk1
        SubChunk2
        ...
      )
      ...
    }
    ...
  )
  [ 'idx1' (<AVI Index>) ]
)
```

Listing A.1: Detailed AVI RIFF form

A.2 The AVI and Bitmapinfo headers

Listing A.2 shows the definition of the main AVI header.

```
typedef struct _avimainheader {
    char        fcc [4];
    int32_t     cb;
    int32_t     dwMicroSecPerFrame;
    int32_t     dwMaxBytesPerSec;
    int32_t     dwPaddingGranularity;
    int32_t     dwFlags;
    int32_t     dwTotalFrames;
    int32_t     dwInitialFrames;
    int32_t     dwStreams;
    int32_t     dwSuggestedBufferSize;
    int32_t     dwWidth;
    int32_t     dwHeight;
    int32_t     dwReserved [4];
} AVIMAINHEADER;
```

Listing A.2: The AVIMAINHEADER structure.

Listing A.3 shows the definition of the BITMAPINFOHEADER.

```
typedef struct tagBITMAPINFOHEADER {
    uint32_t    biSize;
    uint32_t    biWidth;
    uint32_t    biHeight;
    uint16_t    biPlanes;
    uint16_t    biBitCount;
    uint32_t    biCompression;
    uint32_t    biSizeImage;
    uint32_t    biXPelsPerMeter;
    uint32_t    biYPelsPerMeter;
    uint32_t    biClrUsed;
    uint32_t    biClrImportant;
} BITMAPINFOHEADER;
```

Listing A.3: The BITMAPINFOHEADER structure.

B || Stegasis example use

B.1 Stegasis usage information

Listing B.1 shows the information displayed to the user when the program is run with no (or incorrect) arguments supplied.

```

/ ----- | |
| ( ----- | | ----- (-)
\ --- \ | --/ - \/ - \ | / - \ / -- | / -- |
----- ) | || --/ ( - | ( - | \ -- \ \ -- \
| ----- / \ -- \ ----- | \ , \ \ , - | --- / - | --- /
              --/ v3.7 a
              | --- /

```

Stegasis usage:

```
stegasis <command> [-p,-f] --alg=<alg> [--crypt=<alg>] --pass=<pass> [--  
pass2=<pass2>] --cap=<capacity> <video_path> <mount_point>
```

Example usage:

```
stegasis format --alg=lsbp --crypt=aes --pass=password123 --cap=50 /media
/video.avi
stegasis mount --alg=lsbp --crypt=aes --pass=password123 /media/video.avi
/tmp/steg
```

Commands :

format	Formats a video for use with stegasis
mount	Mounts a formatted video to a given mount point

Required Flags:

- alg Embedding algorithm to use, see below
- pass Passphrase used for encrypting and permuting data
- cap Percentage of frame to embed within in percent

Optional flags :

```

—crypt  Cryptographic algorithm used to encrypt embedded data
—pass2  Passphrase used for encrypting and permuting the hidden volume
-p      Do not flush writes to disk until unmount
-f      Force the FFmpeg decoder to be used

```

Embedding Algorithms:

Uncompressed AVI only:

- lsb: Least Significant Bit Sequential Embedding
- lsbp: LSB Permuted Embedding using a seeded LCG

Other video formats:

dctl:	LSB Sequential Embedding within DCT coefficients
dctp:	LSB Permuted Embedding within DCT coefficients
f4:	Implementation of the F4 embedding algorithm
f5:	Implementation of the F5 embedding algorithm

Cryptographic Algorithms:

```

aes:      256 bit AES (Rijndael)
twofish:  256 bit TwoFish
serpent:  256 bit Serpent
aes_serpent_twofish: Chained combination of the above 3

```

Listing B.1: Stegasis usage information.

B.2 Example use case

Listing B.2 shows an example use of **Stegasis**, formatting and mounting a video with complete command line output.

```
$ stegasis format -alg=lsbp -crypt=aes -pass=hunter2 -cap=20 /media/Backup/
video.avi

  -----
 /  ---- | |                                     (-)
 | ( --- | | - --- - - - - - - - - - - - -
 \ --- \| --/ - \| - \ / - \ / - \ / - \| / --|
 ----) | || --/ (- | | (- | \ -- \ \ -- \
 | ----/ \ -- \ --- \| --, \| --, -| ---/-| ---/
                               --/ | v3.7a
                               | ---/
Filesize: 776143108
Totalframes: 280
Width: 1280
Height: 720

Reading AVI chunks...
100% [=====]
Finished parsing AVI file

Volume capacity: 19.32MB

Writing back to disc...
100% [=====]
Format successful!

$ stegasis mount -alg=lsbp -crypt=aes -pass=hunter2 /media/Backup/video.avi
/tmp/steg

  -----
 /  ---- | |                                     (-)
 | ( --- | | - --- - - - - - - - - - - - -
 \ --- \| --/ - \| - \ / - \ / - \ / - \| / --|
 ----) | || --/ (- | | (- | \ -- \ \ -- \
 | ----/ \ -- \ --- \| --, \| --, -| ---/-| ---/
                               --/ | v3.7a
                               | ---/
Filesize: 776143108
Totalframes: 280
Width: 1280
Height: 720

Reading AVI chunks...
100% [=====]
Finished parsing AVI file

Header: STEG
Mounting...

[Second Terminal]
$ cd /tmp/steg
```

```
/tmp/steg$ cp ~/vid.mp4 .

[Stegasis Terminal]
Embedding, Frame: 203, Size: 36158, Offset: 0
Writing back to disc...
100% [=====]

[Second Terminal]
/tmp/steg$ ls -lah
total 4.0K
drwxr-xr-x  2 root root    0 Jan  1  1970 .
drwxrwxrwt 10 root root  4.0K Jan 15 21:59 ..
-rwxr-xr-x  1 root root 14M Jan  1  1970 vid.mp4

[Stegasis Terminal]
<Control-C>
Unmounting...
Writing back to disc...
100% [=====]
Successfully unmounted
```

Listing B.2: Stegasis example use.

C || Detailed Code samples

This section contains select detailed code samples.

C.1 Chi-Squared attack

Listing C.1 shows an implementation of the Chi-Squared attack in Python which operates on the PPM image format.

```
from scipy import integrate, special

file_bytes = open(sys.argv[1], 'rb').read()

header = file_bytes[:2]
if header != 'P6':
    print 'File is not a P6 ppm.'
    sys.exit(0)

fp = 3
width = ''
while True:
    b = file_bytes[fp]
    fp += 1
    if b == ' ':
        break
    width += b

height = ''
while True:
    b = file_bytes[fp]
    fp += 1
    if b == ' ':
        break
    height += b

max_pixel_val = ''
while True:
    b = file_bytes[fp]
    fp += 1
    if b == ' ':
        break
    max_pixel_val += b

output = [0]*100

# fp is now on the first pixel red byte
frameStart = fp
for h in range(1, 100):
    fp = frameStart
    totalPixels = math.floor((h/100.0)*int(width)*int(height))
    X = [0]*(128*3) # X[k] = frequency(2k)
    Y = [0]*(128*3) # Y[k] = frequency(2k+1)
```

```

Z = [0.0]*(128*3)

# Populate the frequency arrays
end = fp + totalPixels*3
while fp < end:
    b = ord(file_bytes[fp])
    if b % 2 == 0:
        X[b/2] += 1
    else:
        Y[(b-1)/2] += 1
    fp += 1

# Calculate theoretically expected frequency
for i in range(len(Z)):
    Z[i] = (X[i] + Y[i]) / 2.0

n = 128
# Minimum frequency condition
for k in range(127):
    if X[k] + Y[k] <= 4:
        X[k] = 0
        Y[k] = 0
        n -= 1

X2 = 0.0
for i in range(128):
    if Z[i] == 0:
        continue
    X2 += ((X[i] - Z[i])**2) / Z[i]

# Calculate probability of embedding
p = 1.0 - special.gammainc((n-1)/2.0, X2/2.0)
output[h] = p

# Print results to stdout
for i in range(1, len(output)):
    print str(i) + " " + str(output[i])

```

Listing C.1: Chi-Squared attack Python implementation (script/pov3.py).

C.2 Reading and writing to the file system

Listing C.2 shows the final implemented version of the FUSE read function callback.

```
int SteganographicFileSystem::read(const char *path, char *buf, size_t size,
    off_t offset, struct fuse_file_info *fi) {
    unordered_map<string, int>::const_iterator file = this->fileSizes.find(
        path);

    if (file == this->fileSizes.end() || offset > file->second)
        return -ENOENT;
    if (size + offset > file->second)
        size = file->second - offset;

    vector<FileChunk> fileChunks = this->fileIndex[path];
    int bytesRead = 0, chunkNum = 0, bytesWritten = 0;

    for (struct FileChunk c : fileChunks) {
        if (bytesRead + c.bytes > offset) {
            break;
        } else {
            bytesRead += c.bytes;
            chunkNum++;
        }
    }
    while (bytesWritten < size) {
        struct FileChunk chunk = fileChunks.at(chunkNum);
        int chunkOffset = offset - bytesRead;
        int bytesLeftInChunk = chunk.bytes - chunkOffset;
        bytesLeftInChunk = min((int)(size - bytesWritten), bytesLeftInChunk);

        printf("\e[1A");
        printf("\e[0K\rExtracting bytes: %d, offset: %d, frame: %d\n",
            bytesLeftInChunk, chunk.offset, chunk.frame);
        if (chunkOffset == 0) {
            this->extract((int *)&chunk.frame, (int *)&chunk.offset,
                bytesLeftInChunk, buf+bytesWritten);
        } else {
            char *temp = (char *)malloc(chunk.bytes * sizeof(char));
            this->extract((int *)&chunk.frame, (int *)&chunk.offset, chunk.bytes,
                temp);
            memcpy(buf + bytesWritten, temp + chunkOffset, bytesLeftInChunk);
            free(temp);
        }
        bytesWritten += bytesLeftInChunk;
        bytesRead = offset;
        chunkNum++;
    }
    return size;
};
```

Listing C.2: FUSE read function callback (fs/stegfs.cc:268).

Listing C.3 shows the final implemented version of the FUSE write function callback.

```
int SteganographicFileSystem::write(const char *path, const char *buf,
    size_t size, off_t offset, struct fuse_file_info *fi) {
    int bytesWritten = 0, nextFrame = 0, nextOffset = 0;
    this->decoder->getNextFrameOffset(&nextFrame, &nextOffset);

    struct FileChunk triple;
    triple.frame = nextFrame;
    triple.offset = nextOffset;
    triple.bytes = 0;

    while (bytesWritten < size) {
        printf("\e[1A");
        printf("\e[0K\rEmbedding, nextFrame: %d, nextOffset: %d, bytesWritten: %d\n",
            nextFrame, nextOffset, bytesWritten);
        pair<int, int> written = this->alg->embed(this->decoder->getFrame(
            nextFrame), (char *) (buf+bytesWritten), size-bytesWritten,
            nextOffset);

        triple.bytes += written.first;
        nextOffset = written.second;
        if (written.second == 0) nextFrame++;
        bytesWritten += written.first;
    }
    this->fileIndex[path].push_back(triple);

    this->decoder->setNextFrameOffset(nextFrame, nextOffset);

    if (offset == 0) {
        this->fileSizes[path] = size;
    } else {
        this->fileSizes[path] += size;
    }
    return size;
};
```

Listing C.3: FUSE write function callback (`fs/stegfs.cc:320`).

C.3 Implementing F5

The F5 steganographic algorithm, developed by Andreas Westfeld^[40], embeds data within the LSBs of DCT coefficients within JPEG images. F5 implements matrix encoding to improve embedding efficiency and is not vulnerable to the Chi-Squared or Histogram attacks. The F5 embedding procedure is summarised in Algorithm C.1.

F5 only embeds data within zero valued DCT coefficients. This can lead to a situation known as *shrinkage* which occurs when a coefficient with absolute value 1 is decremented, becoming zero. The extraction processes cannot distinguish between zero valued coefficients skipped and those which were decremented. Therefore when shrinkage occurs,

Algorithm C.1 The F5 algorithm.

- 1: Estimate the embedding capacity of the carrier medium
 - 2: Determine the parameter k from the message size and embedding capacity.
 - 3: Calculate the code word length $n = 2^k - 1$
 - 4: Embed the secret message with $(1, n, k)$ matrix encoding:
 - 5: Fill a buffer with the LSBs of the next n nonzero DCT coefficients
 - 6: Hash this buffer
 - 7: XOR k bits of the message with the hash value
 - 8: If the output i is nonzero, decrement the absolute value of coefficient at index $i - 1$
 - 9: If shrinkage did not occur, advance the message by k bits.
 - 10: If there is still message data, jump to line 5, else return
-

that k -bit message must be embedded again. To estimate the embedding capacity of the carrier medium taking into account possible shrinkage the following formula is used:

$$\mathcal{C} = DCT - DCT/64 - DCT_0 - DCT_1 + 0.49 \cdot DCT_1$$

where DCT is the total number of DCT coefficients in the image, DCT_0 the total number of zero valued coefficients and DCT_1 the total number of coefficients with absolute value 1. $-0.51 \cdot DCT_1$ is the estimated loss due to shrinkage.

The embedding rate is defined as:

$$\mathcal{R}(k) = \frac{k}{2^k - 1}$$

and is used to compute the optimal value for k . From the estimated capacity \mathcal{C} and the message length m , the optimal embedding rate is calculated as $r = \frac{m}{\mathcal{C}}$ we then seek the maximum integer $k \geq 1$ such that:

$$\mathcal{R}(k) > r$$

The hash function used to hash the LSBs of the nonzero DCT coefficients is the bitwise XOR of each element multiplied by the index:

$$f(\mathbf{a}) = \bigoplus_{i=1}^n a_i \cdot i$$

The index of the coefficient to modify (if any) is calculated by XORing k bits of the message with the hash of the coefficient buffer as follows:

$$i = M_k \oplus f(\text{buffer})$$

If $i \neq 0$ then the $(i - 1)$ th buffer elements coefficients absolute value is decremented by 1.

Extraction follows a similar procedure in which blocks of n^{23} nonzero valued DCT coefficients are hashed to give k message bits. These need to be collected together and returned as a byte array.

Listing C.4 shows a cut down version of the final implementation.

```
virtual pair<int,int> embed(Frame *c, char *data, int reqByteCount, int
    offset) {
    this->crypt->encrypt(data, reqByteCount);
    JBLOCKARRAY frame;
    int row, block, co;
    // Estimate the embedding capacity
    int totalCoefficients = this->dec->getFrameHeight() * this->dec->
        getFrameWidth() * 64;
    int zeroCoefficients = 0, oneCoefficients = 0;
    for (int i = 0; i < totalCoefficients; i++) {
        this->getCoef(lcg.map[i], &row, &block, &co);
        frame = (JBLOCKARRAY)c->getFrameData(row, 1);
        if (frame[0][block][co] == 0) zeroCoefficients++;
        if (frame[0][block][co] == 1 || frame[0][block][co] == -1)
            oneCoefficients++;
    }
    // In bits
    int embeddingCapacity = totalCoefficients - totalCoefficients/64 -
        zeroCoefficients - oneCoefficients - 0.49*oneCoefficients;
    // Force to be a multiple of 8
    embeddingCapacity -= embeddingCapacity % 8;
    if (embeddingCapacity < 8) {
        // No point trying to embed anything in this frame
        int currentFrame, currentFrameOffset;
        this->crypt->decrypt(data, reqByteCount);
        return make_pair(0,0);
    }
    int bitsToEmbed = min(reqByteCount * 8, (int)(embeddingCapacity * (this->
        dec->getCapacity()/100.0)));
    double embeddingRate = (double)bitsToEmbed / (double)embeddingCapacity;
    char k = 1;
    while (true) {
        double rate = (double)k / (pow(2, k) - 1);
        if (rate < embeddingRate) break;
        k++;
    }
    k--;
    int codeWordLength = pow(2, k) - 1;
    int bitsEmbedded = 0;
    while (bitsEmbedded < bitsToEmbed) {
        int oldOffset = offset;
        int *coefficients = this->getNextCoefficientBlock(c, &offset,
```

²³There are some fiddly issues here (which are not mentioned within the original paper on F5) since the extraction procedure must know the value of k in advance along with the number of embedded bits within the image. My solution for this was to embed the value of k and the number of embedded bits within the second JPEG component using the basic JSteg algorithm. The F5 extraction process can then access this information.

```
        codeWordLength);
    if (coefficients == NULL) break;
    int hashOfCoefficients = this->hash(coefficients, codeWordLength);
    int dataBlock = this->getNextDataBlock(data, reqByteCount, k,
        bitsEmbedded);
    int index = hashOfCoefficients ^ dataBlock;
    if (index == 0) {
        // Don't need to do anything
        bitsEmbedded += k;
        continue;
    } else {
        // Need to decrement coefficient at index-1
        index--;
        if (this->decCo(c, oldOffset, codeWordLength, index) == 0) {
            // Shrinkage occurred
            offset = oldOffset;
            continue;
        } else {
            bitsEmbedded += k;
            continue;
        }
    }
    free(coefficients);
}

return make_pair(bitsEmbedded / 8, 0);
};
```

Listing C.4: F5 implementation (steg/f5.cc:86).

D || Testing

D.1 Unit Testing

Unit testing the steganographic and cryptographic algorithms involved mocking a number of the system components, for example the `Frame` and `VideoDecoder` objects. Listing D.1 shows the implemented mock frame used within the test suite.

```
class MockFrame : public Frame {
public:
    MockFrame(void *data): data(data) {};
    long getFrameSize() { return 0; };
    char *getFrameData(int n=0, int c=0) { return (char *)this->data; };
    bool isDirty() { return false; };
    void setDirty() {};

private:
    void *data;
};
```

Listing D.1: MockFrame object (`steg/steg_alg_test.cc:27`).

Listing D.2 shows the test suite which is executed for steganographic and cryptographic algorithm combinations.

```
void runTests() {
    printf("Running tests for %s...\n", this->name.c_str());

    this->AlgorithmShouldCorrectlyExtractEmbeddedData();
    this->AlgorithmShouldNotModifyGivenData();
    this->EmbedShouldReturnZeroBytesEmbeddedIfZeroBytesRequested();
    this->ExtractShouldReturnZeroBytesEmbeddedIfZeroBytesRequested();
    this->EmbedShouldEmbedAsMuchAsPossibleOfTheRequestedData();
    this->EmbedShouldReturnOffsetZeroIfCurrentFrameIsFull();
    this->ExtractShouldReturnOffsetZeroIfCurrentFrameIsFinished();
    this->ExtractShouldReturnNothingIfLessThanAByteIsExtracted();

    printf("All tests complete for %s!\n", this->name.c_str());
    printf("-----\n\n");
};
```

Listing D.2: Stego algorithm test suite (`steg/steg_alg_test.cc:63`).

D.2 Integration Testing

As briefly mentioned in the implementation section, the integration test suite makes use of a number of test archives. These archives are decompressed directly into the volume presented by `Stegasis`. The resulting file system is then traversed and the contents of files checked to make sure they match their original content. This process is automated using several bash scripts as shown in Listing D.3.

```
# Run stegasis
(stegasis format -alg=lsbp -pass=test $1)
(stegasis mount -alg=lsbp -pass=test $1 /tmp/test) &
stegasis_pid=$!
# Wait for the video to mount
sleep 5

# Extract the archives
for f in *.tar; do tar xf $f -C /tmp/test; done

expected_files="dirfile2.txt file1.txt file2.txt testdir"
files=$(ls -C /tmp/test)
if [ "$files" != "$expected_files" ]; then
    echo "ls returned incorrect file list"
    echo "$files"
    cleanExit
fi

expected_file_1="This is a test file."
expected_file_2="This is a different test file."
file1=$(cat /tmp/test/file1.txt)
file2=$(cat /tmp/test/file2.txt)
if [[ "$file1" != "$expected_file_1" || "$file2" != "$expected_file_2" ]];
then
    echo "Contents of file(s) incorrect"
    echo "$file1"
    echo "$file2"
    cleanExit
fi

expected_files_sub="dirfile1.txt"
files_sub=$(ls -C /tmp/test/testdir)
if [ "$files_sub" != "$expected_files_sub" ]; then
    echo "ls returned incorrect file list for sub directory"
    echo "$files_sub"
    cleanExit
fi

expected_file_1_sub="I am in a directory."
file1_sub=$(cat /tmp/test/testdir/dirfile1.txt)
if [[ "$file1_sub" != "$expected_file_1_sub" ]]; then
    echo "Contents of sub directory file(s) incorrect"
    echo "$file1_sub"
    cleanExit
fi

(kill $stegasis_pid)
echo -e "\nAll tests passed :)\n"
```

Listing D.3: The simple integration test suite (test/simple_integration_tests.sh).

Other similar test suites exist which test different parts of the file system functionality, for example the moving of files using `mv` and copying large amounts of data using `cp`.

E || User study results

Listing E.1 shows the the raw data collected from the user study.

```
Raw data: 01000001101111001001110000000000001011011000000100101110110101110
0110101101111011110000001011110010011101110010111000010110001110011100
01000110010111001001001100000100010100111101011011101100010001100000011
11011001101000111001110110111111000111001111110110001100111101110100001
11000110100011001111101111001100010100111101000101000111100001101110011
10111101100101001111100111000010001110000111101000010010110010111011111
01101110110110001110100110000000001101100011100110000001010011100000001
00000101010111100010110101010001001110100011011001110100001111001111010
00100110011010010101101001010101111001010011100111001001001111000110101
01110000000100100010110011010001101011101110011001101011111101011100011
10001101000101000101110100100000011100100101000101111010000011001001110
11010101011110111010011011010010001011010111000010010110100100111101000
11110101010110100001111001010010111110011111010110100100010100100100110
10011001010010101110100110010110101111101110110000001111000110001011010
01001111101001000001110101000110010011110001111101110110110110101010100
01100110001001001110101010110000001000111101101000110110101011011000100
01101111011001111110000110101110110110111011100010001101110111011011010
10100111000101001000111001111000010001110101111001011000010010110011101
00010010010111010110101110000111110011010010011001110110010000011100011
00110100110110111111101001111000010101100101110100001111000110010101100
11010110011101110010101110100011011111100110001001000011010011101000101
01010101010000000100101001001101100000111001100100110011011110110000010
11010111110010001111011101101100110101001001001000011000111101100100011
00011010010111100011101000101010010111000101101111010001000111001011101
110011101011110010001100111101010110110110110101011110110101010011011
10011011101111110010101111010101110001100101011010000101010011111110100
10011001110010000010111110001110100010011101110110111011111011100100011
01000001101011000000110010001101001000110110110100000110101011111100010
1001001110000101111000001000101111001101110000001100001100
```

Total : 2040 , Correct : 1042 , Incorrect : 998
Users: 21

Listing E.1: Data from the user study.

F || Original Project Proposal

COMPUTER SCIENCE PART II PROJECT PROPOSAL

STEGANOGRAPHIC FILE SYSTEMS WITHIN VIDEO FILES

Scott Williams, Christ's College
Originators: Scott Williams

October 20, 2014

PROJECT SUPERVISOR: Daniel Thomas

DIRECTOR OF STUDIES: Professor Ian Leslie

PROJECT OVERSEERS: Professor Peter Robinson, Dr Robert Watson

Introduction and Description of the Work

Steganography is the art of hiding messages within inconspicuous objects - a form of covert communication. Whereas cryptography protects only the content of a message, steganography attempts to conceal the fact that the message even exists. Steganography is particularly useful in countries where encryption is illegal or not suitable, e.g. within the UK, where encryption keys can be forced to be handed over.

There exist many freely available programs which offer message hiding functionality within digital media. However, the majority of these programs operate on single image files and therefore impose a hard limit on the size of message you can embed¹. Many programs also constrain the type of message you can embed to be a simple text string. Video files on the other hand can be several gigabytes in size without arousing suspicion² providing an ideal container for multiple (possibly large) sensitive files. A file system interface would enable users to hide any number of files of any type - just by copying / creating files within the mounted volume. For these reasons, the proposed project focuses on steganographic file systems within video files.

I propose to develop an application which allows a file system to be embedded within a user provided video file. The application will also enable mounting and unmounting of video files with contained file systems. As part of the project I intend to explore a number of steganographic embedding algorithms all of which will be selectable within the final application.

An example use of the final product (henceforward referred to as **Stegasis**):

```
# Prepare an existing video file
$ stegasis format -alg=lsb video.avi

# Using stegasis mount we can directly mount the video file
$ stegasis mount video.avi /mnt/volume

# Create a file inside the file system
$ echo "test" > /mnt/volume/test.txt
# Unmount the file system
$ stegasis umount /mnt/volume
```

After doing some initial research on the topic of steganographic file systems, it seems a suitable approach will be to develop a FileSystem in Userspace using the FUSE package. A similar approach was taken within a paper in which a file system was embedded within multiple JPEG images. For the purposes of this project I'll be focusing on uncompressed raw AVI video files.

¹JPEG images for example are typically only a few megabytes in size - limiting the size of files you can possibly embed.

²Raw uncompressed AVI files are roughly 2GB per minute of footage.

I propose a staged approach to the project where each stage implements an increasingly secure scheme of embedding the file system, for example starting with naïve least significant bit embedding, showing how this can be broken using statistical analysis and then moving on to more advanced techniques (each method selectable via the `-alg` flag). The main product of this project - **Stegasis** - as shown above will be a user facing application, enabling all versions of the algorithms described throughout the stages of the project to be run on user provided video files. A number of programs to analyse and break insecure schemes proposed early on during the project could also be produced. This project would tie in nicely with the Part II courses Information Theory and Coding, Digital Signal Processing, and possibly Security II.

Resources Required

I will be using the **C++** programming language to develop **Stegasis** of which I have a good amount of experience with. The virtual File System aspect will be implemented using the **FUSE** package. A scripting language such as **Python** or **MATLAB** may also be used to develop some of the steganalysis tools. Raw **AVI** Video files for testing purposes can be created using **VirtualDub**'s video conversion tools.

I intend to implement the project on my own desktop computer (running Ubuntu 14.04.1 as well as Windows 7) due to convenience and accessibility. However, there is no reason why development could not happen on the PWF machines, should this be needed. Backups will be taken at regular intervals and **Git**, a revision control system will be used (in conjunction with **GitHub**) to preserve multiple versions of the project stored both locally and in an offsite location.

Starting Point

Steganography shares a number of concepts with Cryptography for which an introductory course (Security I) was given last year. I have read the introductory chapters of Steganography in Digital Media by Jessica Fridrich, a number of generic steganography papers and also a few papers specific to Steganographic File Systems.

I have implemented a simple “hello world” **FUSE** virtual file system in **C++** to prove the package works as I would expect.

Substance and Structure of the Project

The project will consist of the following sections:

1. Research and investigation into the theoretical aspects of steganography, identifying appropriate embedding algorithms and steganalysis techniques. Investigation into

developing a virtual file system and the AVI video format.

2. Design and implementation of **Stegasis** providing a variety of steganographic embedding algorithms, allowing raw AVI files to be formatted, and the mounted file system to be written to and read from. This section will follow an iterative process wherein each iteration will propose an increasingly secure embedding algorithm and an attempt to develop a suitable steganalysis technique to break it.
3. Evaluation of **Stegasis** will be based on the following criteria:
 - **Correctness:** **Stegasis** correctly formats and mounts a provided video file presenting the file system as a logical volume. Files written to the volume should persist between unmounts and subsequent mounts of the same unmodified video file.
 - **Usability:** The **Stegasis** command line tool should be simple and intuitive providing usage details for its functionality and helpful error messages.
 - **Performance:** The steganographic embedding process should have no noticeable impact on the file system performance i.e. writes to files should not be perceivably slower than a standard HDD³.

Success Criteria

For the project to be considered a success, **Stegasis** should provide the following functionality:

- **Stegasis** should offer a number of steganographic embedding algorithms.
- Given a standard raw AVI video file, **Stegasis** should format⁴ the video such that it can be mounted.
- Given a formatted video file, **Stegasis** should be able to mount the video and present a virtual file system at a given mount point.
- Standard file system operations including listing files, reading a file, writing to a file and deleting a file should be supported within the virtual file system.
- The above described functionality of **Stegasis** should operate without noticeable visual impact on the video content.

³For example, it takes roughly 2ms to read a 1MB file from a HDD.

⁴Format in this case is referring to writing some meta data to the video e.g. which embedding algorithm is being used.

Extensions

If there is sufficient time, the following extensions may be attempted:

- **Directory Structure:** **Stegasis** as described only permits files to be created within the root of the virtual file system. It would be beneficial to allow users to create folder structures as you would expect from a standard file system.
- **Audio Usage:** **Stegasis** as described only makes use of video image frames to embed the file system. However, a substantial part of an AVI file may be the audio data. It would be useful to make use of the audio data to increase the steganographic capacity of the proposed embedding algorithms.
- **Video Formats:** Unfortunately, raw AVI video is (very) uncommon compared with modern compressed video formats such as H.264 (mp4) it would be very beneficial for **Stegasis** to operate on a variety of video formats, rather than just raw AVI. However, modern video formats are intricate and complex so this may well be outside the scope of this project.
- **Evaluation of video artifacts:** Quantification of “noticeable visual impact on video content” via a set of human trials possibly achievable by crowd-sourcing through an online website.

Timetable

Michaelmas Term

- **24th October - 6th November (weeks 3-4):** Research on the theoretical background of steganography including reading relevant sections of textbooks and academic papers.
- **7th November - 20th November (weeks 5-6):** Investigation of appropriate steganographic embedding algorithms suitable for video files. Investigation of the AVI video format and the FUSE package.
- **21st November - 4th December (weeks 7-8):** Implementation of **Stegasis** only offering the simple LSB embedding algorithm and file system functionality.

Winter Vacation

- **5th December - 18th December:** Implementation of more advanced steganographic embedding techniques and integration of these into **Stegasis**. Development of steganalysis tools to break proposed embedding schemes.
- **19th December - 1st January:** Christmas holiday.

-
- ***2nd January - 16th January:*** Continuing work on more advanced steganographic embedding and steganalysis techniques.

Lent Term

- ***16th January - 22nd January (week 1):*** Polishing of **Stegasis** and source code (not yet including extension work) - core project should be finished at this point. Write progress report and prepare for the progress presentation.
- ***23rd January - 12th February (weeks 2-4):*** Evaluation of the core project. Identification and implementation of promising extension tasks.
- ***13th February - 26th March (weeks 5-10):*** Work on the dissertation write-up, completing a draft for submission to my supervisor.
- ***27th March - 23rd April (Easter Vacation):*** Revision of dissertation addressing supervisors comments. Dissertation should be ready to submit by 23rd April.